

Gradient-Aware Adaptive Rounding for Post-Training Quantization of Deep Neural Networks

48246347 Piye Liu

Supervisor: Prof. Hiroshi Nakamura

The recent exponential scaling of Large Language Models (LLMs) has resulted in massive carbon emissions and significant challenges in deployment, forcing users to rely on third-party cloud providers, thereby raising privacy concerns. This highlights the necessity of model compression, and post-training quantization (PTQ) is one method that enables efficient on-device models. AdaRound, an adaptive rounding method that outperforms round-to-nearest rounding during quantization, minimizes the task loss between the original weights and the quantized weights on unlabelled calibration data. We find that: 1) AdaRound leaks some information about its calibration data into the weight perturbation matrix, and we demonstrate this by proposing a novel inversion attack that targets the null space of this matrix to recover coarse representations of the original data. 2) The theoretical foundation of AdaRound asserts that the gradient of the loss function is zero for pretrained converged models, an assumption which is refuted by recent literature. To address this discrepancy, we attempted to reconcile this by incorporating local gradient signals into the AdaRound optimization process, but obtained neutral results on empirical evaluations. Therefore, this thesis highlights a frequently overlooked dimension in PTQ (privacy), and proposes a gradient-aware modification of AdaRound to better align the method with modern literature.

Outline

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Exponential Scaling of LLMs

Recently, the exponential scaling of LLM sizes has inadvertently resulted in models that are extremely challenging to deploy, thus causing users to rely on third-party providers for LLM usage.

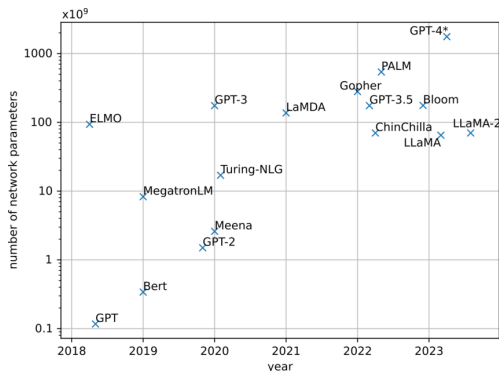


Figure 1: Observed exponential scaling of LLM sizes over time [1] (Note: The GPT-4 figure is an estimate)

Problems with Third-Party APIs

- ▶ **Environmental Effects:** GPT-3's training had an estimated CO₂e of ~552.1 tons [2], GPT-4's training had an estimated CO₂e of ~6912 tons [3], and Llama 3.3 [4]'s training had an estimated CO₂e of ~11,390 tons.
- ▶ **Cost:** For example, OpenAI o1 [5] in “pro mode” costs \$150 per 1 million tokens input, and \$600 per 1 million tokens output. Uber's yearly AI budget was exhausted in 4 months, while Microsoft began cancelling Claude Code licenses just 6 months after acquiring them [6].
- ▶ **Reliability:** Single point of failure, and an outage can have severe effects on mission-critical tasks. For example, it is observed that whenever ChatGPT has an outage, stock trading volume goes down [7].

Benefits of On-Device Inference

- ▶ **Privacy:** Data does not need to leave the device
- ▶ **Robustness:** Can function even without internet access
- ▶ **Latency:** Doesn't have to rely on third-party scheduling, can prioritize urgent tasks
- ▶ **Cost:** Reduces API fees, network costs, and even environmental emissions

Therefore, our goal is to improve model compression to enable efficient on-device AI. Furthermore, we also explore a side effect of model compression (privacy): data-dependent post-training quantization may leak some calibration data.

Model Compression Methods

Methods typically used for the compression of deep neural networks:

- ▶ Weight Pruning: Setting some weights to zeroes, resulting in sparse weight matrices [8].
- ▶ Knowledge Distillation: Smaller model learns to match output logits of larger model [9].
- ▶ Weight Palettization: Clusters weights (e.g. via k -means), then stores every weight as the centroid of the cluster they belong to [10].
- ▶ Quantization: **We focus on this method for our thesis.** Stores weights in a data type of less bit width (e.g. int8), compared to the original data type (e.g. float32) [11].

Objective & Contributions

Overall, our goal is to investigate model compression and its side effects (specifically, privacy) to enable efficient on-device AI. We achieve this by the following.

- ▶ A novel inversion method against AdaRound to recover coarse outlines that represent its calibration data.
- ▶ A proposed gradient-aware adaptive rounding method based on AdaRound, applied on two language models.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Compressibility of Deep Neural Networks

Source coding theorem [12]:

- ▶ n i.i.d. (independent and identically distributed) random variables, each with entropy $H(X)$, can be compressed to $nH(X)$ bits on average, with negligible information loss, as $n \rightarrow \infty$

In this case, $H(X)$, the entropy of a random variable, is the expected amount of information needed to describe it, and can be computed by:

$$H(X) := - \sum_{x \in X} p(x) \log_2 p(x). \quad (1)$$

Note that the log base determines the units: $\log_2$ yields bits, $\log_{10}$ yields dits, and $\ln$ yields nats.

Measuring Compressibility of LLMs

Total information stored in LLM weights: $\sim n \cdot \hat{H}(W)$ bits;
Maximum possible total entropy of the model: $n \cdot b$ bits. It therefore follows that the capacity utilization by a model is:

$$U(W) := \frac{\hat{H}(W)}{b}. \quad (2)$$

We use this method to quantify the capacity utilization of several modern LLMs:

Table 1: Entropy Estimate of Language Models (bfloat16)

Model	$\hat{H}(W)$ in bits	$U(W)$
Qwen 3 4B [13]	10.58	$10.58/16 \approx 66\%$
Qwen 3.5 4B [14]	10.56	$10.56/16 \approx 66\%$
Gemma 4 12B Instruct [15]	10.69	$10.69/16 \approx 67\%$
GLM 4.7 Flash [16] (30B)	10.49	$10.49/16 \approx 66\%$

Small Language Models

- ▶ TinyStories [17] proposed a synthetic low-vocabulary dataset and several small language models, as low as 1M parameters, that produce coherent English.
- ▶ **Limitations:** their usage of the default GPT-Neo [18] tokenizer (which is identical to the pretrained GPT-2 [19] tokenizer), restricted to the top 10,000 tokens.
- ▶ **Our models:** we train two small language models on TinyStories v1 and TinyStories v2, respectively. We use a SentencePiece Byte-Pair Encoding (BPE) tokenizer [20] with a vocabulary size of 1024, and we used a LLaMa-based [21] model with roughly half the parameters of the smallest reference model.

Evaluation of Our LMs

We use perplexity [22] as a quantitative metric to evaluate a language model's accuracy. For a model θ , a text x of length N , it is defined as the following.

$$\text{Perplexity}(\theta, x) = \left(\prod_{i=1}^N \frac{1}{p_{\theta}(x_i \mid x_1, x_2, \dots, x_{i-1})} \right)^{\frac{1}{N}} \quad (3)$$

In other words, it is the geometric mean of the inverse probabilities of a given text.

Table 2: Evaluation of Our Language Models

	Reference	Ours (TinyStories v1)	Ours (TinyStories v2)
Architecture	GPT-Neo	LLaMa	LLaMa
Parameter Count	1.17M	0.57M	0.46M
Training Data	TinyStories v1		TinyStories v2
Tokenizer (Vocab Size)	GPT-Neo	SentencePiece	SentencePiece
Vocab Size	10000	1024	1024
Perplexity $\downarrow$ (TinyStories v1)	19.3919	11.9006	—
Perplexity $\downarrow$ (TinyStories v2)	—	—	8.0994

Baseline Round-to-Nearest Quantization

Consider the symmetric uniform quantization method to int8:

- ▶ Quantization:

- ▶ $\text{scale} = \frac{\max(|W|)}{2^7 - 1}$

- ▶ $W_{\text{quantized}} = \text{clamp}(\lfloor \frac{W}{\text{scale}} \rceil)$

- ▶ Reconstruction (dequantization):

- ▶ $W_{\text{reconstructed}} = \text{scale} \cdot W_{\text{quantized}}$

- ▶ Note that $\text{clamp}()$ here refers to clamping a number between $[-2^7, 2^7 - 1]$ to ensure its within the range of the data type.

Evidently, model weights are perturbed in the direction of $W_{\text{reconstructed}} - W$ after quantization and dequantization. We investigate the effects of perturbation direction on model accuracy on our TinyStories V2 model.

Loss Landscape: Observations

We explored the loss landscape of our TinyStories model and found:

- ▶ When moving along a line, there was a “basin” of low loss, and moving towards the negative direction caused less degradation than moving towards the positive direction.
- ▶ There was a huge outlier of high perplexity when moving randomly.
- ▶ Moving towards another model resulted in less degradation despite the massive distance.

We therefore find that: **When weights of a model are perturbed, the loss does not degrade uniformly.** Some directions result in massive degradation while some directions preserved more accuracy.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

AdaRound: Objective

AdaRound [23] is a novel rounding method, which aims to minimize the difference in task loss of the quantized model and the unquantized model. The task loss of layer j of the unquantized model is $\mathcal{L}(x, y, w_j)$, and the task loss of the same layer of the quantized model $\mathcal{L}(x, y, w_j + \Delta w_j)$. Therefore, the objective is:

$$\arg \min_{\Delta w_j} (\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j)). \quad (4)$$

Using second order Taylor expansion:

$$\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j) \approx \nabla \mathcal{L}(x, y, w_j)^T \Delta w_j + \frac{1}{2} \Delta w_j^T \mathbf{H} \Delta w_j \quad (5)$$

AdaRound: Approximations

The authors assert that $\nabla \mathcal{L}(x, y, w_j)$ is negligible, which transforms the objective into just:

$$\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j) \approx \frac{1}{2} \Delta w_j^T \mathbf{H} \Delta w_j \quad (6)$$

Finally, the authors formulate the optimization of weight rounding for layer j as the following Quadratic Unconstrained Binary Optimization (QUBO) problem.

$$\arg \min_{\Delta w_j} \Delta w_j^T \mathbf{H} \Delta w_j \quad (7)$$

The above problem is NP-hard, so AdaRound uses the local mean squared error alongside a regularizer as a proxy.

$$\arg \min_{\mathbf{V}} \left(\|\mathbf{W}_X - \hat{\mathbf{W}}_X\|_F^2 + f_{\text{reg}}(\mathbf{V}) \right) \quad (8)$$

AdaRound: Algorithm

Algorithm 1 AdaRound

Require: $\mathbf{W}$, X (calibration data), t (number of iterations), b (target bit width), β (regularization parameter)

Ensure: Quantized model $\mathbf{W}_Q$, scales

- 1: $x_{in} \leftarrow \text{CACHEACTIVATIONS}(\mathbf{W}, X)$ ▷ Forward pass on calibration data to cache input activations
 - 2: **for** w_j in $\mathbf{W}$ **do**
 - 3: $q_{\min} \leftarrow -2^{b-1}$
 - 4: $q_{\max} \leftarrow 2^{b-1} - 1$
 - 5: $\text{scales}_j \leftarrow \frac{\max(w_j)}{q_{\max}}$ ▷ Store scale for output
 - 6: $\mathbf{V}_j \leftarrow h^{-1}(w_j - \lfloor w_j \rfloor)$ ▷ Initialize $\mathbf{V}_j$
 - 7: **for** $i \in \{0 \dots t\}$ **do**
 - 8: $\beta \leftarrow \text{ANNEAL}(\beta)$
 - 9: $\hat{w}_j \leftarrow \text{clamp}\left(\lfloor \frac{w_j}{\text{scales}_j} \rfloor + h(\mathbf{V}_j), q_{\min}, q_{\max}\right)$
 - 10: $\mathcal{L} \leftarrow \|w_j x_{in} - \hat{w}_j x_{in}\|_F^2 + \sum_{i,j} (1 - |2h(\mathbf{V}_{i,j}) - 1|^\beta)$
 - 11: $\mathbf{V} \leftarrow \mathbf{V} - \nabla \mathcal{L}$ ▷ via autograd
 - 12: **end for**
 - 13: $\mathbf{W}_{Q_{i,j}} \leftarrow \begin{cases} \lceil w_{i,j} \rceil, & \text{where } h(\mathbf{V}_{i,j}) \geq 0.5 \\ \lfloor w_{i,j} \rfloor, & \text{otherwise} \end{cases}$ ▷ Finalize quantization
 - 14: **end for**
 - 15: **return** $\mathbf{W}_Q$, scales
-

Limitations of AdaRound

- ▶ **Privacy** of calibration data: AdaRound can cause some information to be encoded into the weight perturbation matrix. We empirically verify this in Contribution 1.
- ▶ The assumption in Equation 5 that the gradient of the loss function for pretrained converged models is negligible (for which empirical evidence was not provided) is **challenged by recent literature**.

AdaRound's Zero-Gradient Assumption

While one might expect gradient descent to follow a smooth, stable path to an optimum, in reality it almost always operate in a highly dynamic, oscillatory regime described as the **Edge of Stability** [24], leading to chaotic loss curves such as:




Figure 2: Training Loss Curve of Kimi K2 [25]

The oscillatory optimization also implies that the gradient for pretrained models are rarely zero. Recent literature [24] empirically show that even as the training loss decreases, the gradient norm of a model almost stays constant.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Leakage Channel: The Null Space

- ▶ AdaRound's objective is to make $\|Wx - \hat{W}x\|_F^2$ approach zero. Let $\Delta W = W - \hat{W}$. For a model quantized using AdaRound:

$$\Delta Wx \approx 0$$

- ▶ So the calibration data x is approximately in the **null space** of the perturbation matrix. In other words, $x \in \mathcal{N}(\Delta W)$. We can therefore exploit this property to generate samples that resemble the mean of the calibration data.
- ▶ We **can't** reconstruct exact pixel-for-pixel calibration samples. However, we demonstrate that by implementing this as a simple optimization problem, we can reconstruct representations that look like the mean of these calibration samples.

Realizing an Attack

- ▶ If an attacker has access to both the original weights (W) and the quantized weights ($\hat{W}$), they can recover representations of the calibration data by searching for any vector that satisfies this null space property.
- ▶ For example, hypothetically: a company downloads an open-source vision model, quantizes it with AdaRound using confidential internal data (e.g. patient records), then deploys it onto edge devices sold to customers. An attacker can reconstruct representations of the calibration data by extracting the quantized $\hat{W}$ from the edge devices, then performing the inversion on the weight perturbation matrix $\Delta W = W - \hat{W}$.

Proposed Inversion Algorithm

Algorithm 2 AdaRound Inversion Algorithm

Require: W , $\hat{W}$, $\hat{x}_{\text{init}}$, No. of Iterations K , Learning Rate η

1: $\Delta W \leftarrow W - \hat{W}$

2: $\hat{x} \leftarrow \hat{x}_{\text{init}}$

3: **for** $k = 1$ **to** K **do**

4: $\mathcal{L} \leftarrow \sum (\Delta W \tanh(\hat{x}))^2$ (Forward pass)

5: $\frac{\partial \mathcal{L}}{\partial \hat{x}} \leftarrow \text{Autograd}(\mathcal{L}, \hat{x})$ (Backward pass)

6: $\hat{x} \leftarrow \hat{x} - \eta \frac{\partial \mathcal{L}}{\partial \hat{x}}$ (Update $\hat{x}$)

7: **end for**

8: **return** $\tanh(\hat{x})$

Simple Optimization

Note that the objective can be rewritten as:

$$\|\Delta W \hat{x}\|_2^2 = (\Delta W \hat{x})^T (\Delta W \hat{x}) = \hat{x}^T (\Delta W^T \Delta W) \hat{x}. \quad (9)$$

We can therefore make this simple optimization:

Algorithm 3 Optimized AdaRound Inversion Algorithm

Require: $W, \hat{W}, \hat{x}_{\text{init}},$ No. of Iterations $K,$ Learning Rate η

- 1: $\Delta W \leftarrow W - \hat{W}$
 - 2: $\hat{x} \leftarrow \hat{x}_{\text{init}}$
 - 3: $G \leftarrow \Delta W^T \Delta W$
 - 4: **for** $k = 1$ **to** K **do**
 - 5: $\mathcal{L} \leftarrow \sum (\tanh(\hat{x})^T G \tanh(\hat{x}))$ (Forward pass)
 - 6: $\frac{\partial \mathcal{L}}{\partial \hat{x}} \leftarrow \text{Autograd}(\mathcal{L}, \hat{x})$ (Backward pass)
 - 7: $\hat{x} \leftarrow \hat{x} - \eta \frac{\partial \mathcal{L}}{\partial \hat{x}}$ (Update $\hat{x}$)
 - 8: **end for**
 - 9: **return** $\tanh(\hat{x})$
-

Experiment Setup

- ▶ We attempt an inversion on MLPs trained on the Olivetti faces dataset [26], the Weizmann Horse Database[27], and a Chest X-Ray Images dataset[28] using an MLP classifier quantized by AdaRound. Note that this method can also be applied to LLMs, but we are using image models as a proof-of-concept.
- ▶ For each model, we performed 500 iterations of AdaRound quantization. We then attempted to invert the process by training 10 random Gaussian samples for 3,000 iterations using the Adam optimizer. We report SSIM, PSNR, and Normalized RMSE before and after the inversion.

Quantitative Results


Table 3: Metrics from Random Initialization vs. Our Method

	SSIM $\uparrow$	PSNR $\uparrow$	NRMSE $\downarrow$
Olivetti Faces (Before)	0.00	3.50	1.16
Olivetti Faces (After)	0.14 (+0.14)	12.83 (+9.33)	0.40 (-0.76)
Chest X-Ray (Before)	0.00	2.61	1.10
Chest X-Ray (After)	0.07 (+0.07)	10.13 (+7.52)	0.46 (-0.64)
Horse DB (Before)	0.00	2.60	1.11
Horse DB (After)	0.07 (+0.07)	10.15 (+7.55)	0.47 (-0.64)


The increase in PSNR and the decrease in RMSE across all datasets (mean over 10 samples) indicate that our method is recovering real structural information, although it is very coarse.

Qualitative Results

Let us evaluate our results qualitatively by looking at the generated images. We observe an outline of a face, and we observe the structure of the lungs and the spine.



(a) Results on Olivetti Faces



(b) Results on Chest X-Ray Images (Pneumonia)

Figure 3: Qualitative Results of Inversion Algorithm

Qualitative Results Continued

On the Weizmann horse database, we can see a coarse silhouette of a horse.




Figure 4: Results on Weizmann Horse Database

Ablation Study

- ▶ We perform the same inversion method against the same models quantized with round-to-nearest (RTN) rounding.
- ▶ Since RTN does not use calibration data, we compare our method's samples to the mean of the training data instead.

Table 4: Results from Ablation Study

	SSIM $\uparrow$	PSNR $\uparrow$	NRMSE $\downarrow$
Olivetti Faces (Before)	0.00	2.89	1.12
Olivetti Faces (After)	0.00	3.90 (+1.01)	1.00 (-0.12)
Chest X-Ray (Before)	0.00	2.77	1.11
Chest X-Ray (After)	0.00	3.64 (+0.87)	1.00 (-0.11)
Horse DB (Before)	0.00	2.57	1.10
Horse DB (After)	0.00	3.38 (+0.81)	1.00 (-0.10)

- ▶ We notice 0.00 SSIM on all datasets, implying that the generated samples do not look like the calibration data at all.

Discussion & Limitations

- ▶ **Evidence of Leakage:** Even though the SSIM in the main experiment is quite low, we observe empirical evidence of the leakage channel. Furthermore, the SSIM in the ablation study is 0.00 on all datasets, which empirically shows that AdaRound is likely to be the source of the leakage channel.
- ▶ **Main Limitation:** Our method only reconstructs samples that look like the mean of the training data. It does not reconstruct individual samples as-is.
- ▶ Furthermore, the SSIM we obtained is rather low and there is much room for further improvement.
- ▶ Another limitation is that to perform this attack, the weight perturbation matrix is required. Typically this means needing to obtain both the unquantized and quantized models.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Method

- ▶ To reconcile AdaRound's assumption that the gradient term in Equation 5 can be ignored, and Central Flows' evidence that the gradient is nonzero even for pretrained converged models, we propose a modified version of AdaRound to incorporate awareness of the local gradient.
- ▶ In order to minimize

$$\mathbb{E} \left[g^{(w)^T} \Delta W \right], \quad (10)$$

we can attempt to force ΔW to have the same sign as $-g^{(w)^T}$.

- ▶ We do this by finetuning the model on a small amount of data then taking the EMA of the model during finetuning as W'_{avg} . We add the Binary Cross Entropy (BCE) between the current rounding directions and $\text{sign}(W'_{\text{avg}} - W)$ as a penalty term to AdaRound's optimization objective.

Algorithm for EMA Finetuning

Algorithm 4 Gradient-Aware Adaptive Rounding (Finetuning)

Require: $\mathbf{W}$, X (calibration data), η (learning rate), B (batch size)

Ensure: EMA of models during finetuning $\mathbf{W}'_{\text{avg}}$

- 1: $\mathbf{W}' \leftarrow \mathbf{W}$
 - 2: $\mathbf{W}'_{\text{avg}} \leftarrow \mathbf{W}$
 - 3: $N \leftarrow \frac{X}{B}$ ▷ Number of Iterations
 - 4: $\alpha \leftarrow \frac{2}{N+1}$ ▷ EMA Smoothing Factor
 - 5: **for** each batch x_b in X **do**
 - 6: $\hat{x}_b \leftarrow \mathbf{W}'(x_b)$ ▷ Forward pass
 - 7: $\mathcal{L} \leftarrow \text{CrossEntropyLoss}(\hat{x}_b, x_b)$
 - 8: $\mathbf{W}' \leftarrow \mathbf{W}' - \eta \cdot \nabla \mathcal{L}$ ▷ Backward pass
 - 9: $\mathbf{W}'_{\text{avg}} \leftarrow \alpha \cdot \mathbf{W}'_{\text{avg}} + (1 - \alpha) \cdot \mathbf{W}'$ ▷ Update EMA
 - 10: **end for**
 - 11: **return** $\mathbf{W}'_{\text{avg}}$
-

Proposed Method Objective

Our proposed method is essentially a modification to AdaRound's objective in Equation 8:

$$\arg \min_{\mathbf{V}} \left(\|\mathbf{W}_X - \hat{\mathbf{W}}_X\|_F^2 + f_{\text{reg}}(\mathbf{V}) + \underbrace{\psi \cdot \text{BCE} \left(h(\mathbf{V}_j), \text{sign}(\mathbf{W}' - \mathbf{W}) \right)}_{\text{Proposed addition}} \right) \quad (11)$$

Experiment Setup

We evaluate the proposed method on our TinyStories models with the following hyperparameters and components.

Table 5: Implementation Details

Hyperparameter/Component	Value
Finetuning Data	2% of Training Data
Learning Rate	3e-4
Optimizer	AdamW [29]
Learning Rate Scheduler	CosineAnnealingLR [30]
ψ	100.0

Results: Main Experiment

- ▶ We observe **neutral to negative** results using the proposed method.
- ▶ Note that the two models had different testing data (TinyStories v1 and TinyStories v2, respectively).

Table 6: Perplexity Evaluation of Proposed Method

Type of Quantization	TinyStories V1	TinyStories V2
None (Unquantized Model)	11.9006	8.0994
Round-to-Nearest	11.9219	8.1224
AdaRound	11.9216	8.1177
Proposed Method	11.9216	8.1179

Ablation Study

We report an ablation study by ablating two possible elements of the proposed method: AdaRound and EMA. Ablating EMA means we use final finetuned model weights instead of the EMA of model weights during all iterations.

Table 7: Perplexity Evaluation of Ablation Studies

	AdaRound+BCE	EMA	TinyStories V1	TinyStories V2
Main Experiment	✓	✓	11.9216	8.1179
Ablation 1	✓	×	11.9216	8.1197
Ablation 2	×	✓	11.9166	8.1546
Ablation 3	×	×	11.9050	8.1084

Note again that the two models had different testing data (TinyStories v1 and TinyStories v2, respectively).

Ablation Study: Method

Note that ablating AdaRound means we steer the rounding using only the sign of the weight perturbation matrix:

Algorithm 5 Gradient-based Steering without AdaRound

Require: $\mathbf{W}$ (unquantized weights), $\mathbf{W}'$ (calibration model), b (target bit width)

Ensure: Quantized model $\mathbf{W}_q$

1: **function** CUSTOMROUND($\mathbf{W}$, $\mathbf{W}'$)

2: $\mathbf{W}_{i,j} \leftarrow \begin{cases} \lceil \mathbf{W}_{i,j} \rceil, & \text{where } \mathbf{W}_{i,j} < \mathbf{W}'_{i,j} \\ \lfloor \mathbf{W}_{i,j} \rfloor, & \text{otherwise} \end{cases}$

3: **return** $\mathbf{W}$

4: **end function**

5: $q_{\min} \leftarrow -2^{b-1}$

6: $q_{\max} \leftarrow 2^{b-1} - 1$

7: $\text{scale} \leftarrow \frac{\max(\mathbf{W})}{q_{\max}}$

8: $\mathbf{W}_q \leftarrow \text{clamp}\left(\text{CUSTOMROUND}\left(\frac{\mathbf{W}}{\text{scale}}\right), q_{\min}, q_{\max}\right)$

9: **return** $\mathbf{W}_q$, scale

Discussion

- ▶ We notice that we achieved the best performance by ablating both EMA and AdaRound (outperforming original AdaRound)
- ▶ This ablation study shows that the gradient direction can still potentially be useful despite the neutral to negative results in the main experiment.
- ▶ However, it is also possible that while the gradient is non-zero, the assertion that it is zero is simply “good enough” for the search of the optimal rounding direction.
- ▶ Finally, it might be possible to consider a better proxy than the direction towards a finetuned model.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Conclusion

- ▶ **AdaRound Inversion:** We found that AdaRound inadvertently encodes information about its calibration data into its weight perturbation matrix. We propose a novel inversion method to reconstruct representations of the calibration data and empirically confirmed the existence of the leakage channel.
- ▶ **Gradient-Aware AdaRound:** We challenged an assertion in AdaRound that the gradient of the loss function for pretrained converged models is negligible. However, we obtained neutral to negative results using our proposed method. An ablation study actually outperformed AdaRound, potentially implying that the gradient for pretrained converged can still be a useful signal.

Future Work

In the future, we would like to:

- ▶ Improve the quality of the inversion attack using regularizers (e.g. Total Variation (TV) loss) and priors (e.g. deep image priors).
- ▶ Furthermore explore the incorporation of the gradient of the loss function in AdaRound, possibly using other proxies to estimate the direction of the gradient.
- ▶ Replicate our experiments on models of larger scale, so we can determine whether its behavior persists beyond smaller on-device models.

References

- [1] J. Gerstmayr, P. Manzl, and M. Pieber, "Multibody models generated from natural language," *Multibody System Dynamics*, vol. 62, pp. 249–271, Jan. 2024.
- [2] D. Patterson, J. Gonzalez, Q. Le, C. Liang, L.-M. Munguia, D. Rothchild, D. So, M. Texier, and J. Dean, "Carbon Emissions and Large Neural Network Training," Apr. 2021.
- [3] Z. Kausar, S. Latif, R. K. Shahzad, and M. Fatima, "Quantifying the Climate Risk of Generative AI: Region-Aware Carbon Accounting with G-TRACE and the AI Sustainability Pyramid," May 2026.
- [4] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, A. Yang, A. Fan, A. Goyal, A. Hartshorn, A. Yang, A. Mitra, A. Sravankumar, A. Korenev, A. Hinsvark, A. Rao, A. Zhang, A. Rodriguez, A. Gregerson, A. Spataru, B. Roziere, B. Biron, B. Tang, B. Chern, C. Caucheteux, C. Nayak, C. Bi, C. Marra, C. McConnell, C. Keller, C. Touret, C. Wu, C. Wong, C. C. Ferrer, C. Nikolaidis, D. Allonsius, D. Song, D. Pintz, D. Livshits, D. Wyatt, D. Esiobu, D. Choudhary, D. Mahajan, D. Garcia-Olano, D. Perino, D. Hupkes, E. Lakomkin, E. AlBadawy, E. Lobanova, E. Dinan, E. M. Smith, F. Radenovic, F. Guzmán, F. Zhang, G. Synnaeve, G. Lee, G. L. Anderson, G. Thattai, G. Nail, G. Mialon, G. Pang, G. Cucurell, H. Nguyen, H. Korevaar, H. Xu, H. Touvron, I. Zarov, I. A. Ibarra, I. Kloumann, I. Misra, I. Evtimov, J. Zhang, J. Copet, J. Lee, J. Geffert, J. Vranes, J. Park, J. Mahadeokar, J. Shah, J. van der Linde, J. Billock, J. Hong, J. Lee, J. Fu, J. Chi, J. Huang, J. Liu, J. Wang, J. Yu, J. Bitton, J. Spisak, J. Park, J. Rocca, J. Johnstun, J. Saxe, J. Jia, K. V. Alwala, K. Prasad, K. Upasani, K. Plawiak, K. Li, K. Heafield, K. Stone, K. El-Arini, K. Iyer, K. Malik, K. Chiu, K. Bhalla, K. Lakhotia, L. Rantala-Young, L. van der Maaten, L. Chen, L. Tan, L. Jenkins, L. Martin, L. Madaan, L. Malo, L. Blecher, L. Landzaat, L. de Oliveira, M. Muzzi, M. Pasupuleti, M. Singh, M. Paluri, M. Kardas, M. Tsimpoukelli, M. Oldham, M. Rita, M. Pavlova, M. Kambadur, M. Lewis, M. Si, M. K. Singh, M. Hassan, N. Goyal, N. Torabi, N. Bashlykov, N. Bogoychev, N. Chatterji, N. Zhang, O. Duchenne, O. Çelebi, P. Alrassy, P. Zhang, P. Li, P. Vasic, P. Weng, P. Bhargava, P. Dubal, P. Krishnan, P. S. Koura, P. Xu, Q. He, Q. Dong, R. Srinivasan, R. Ganapathy, R. Calderer, R. S. Cabral, R. Stojnic, R. Raileanu, R. Maheswari, R. Girdhar, R. Patel, R. Sauvestre, R. Polidoro, R. Sumbaly, R. Taylor, R. Silva, R. Hou, R. Wang, S. Hosseini, S. Chennabasappa, S. Singh, S. Bell, S. S. Kim, S. Edunov, S. Nie, S. Narang, S. R. Parthi, S. Shen, S. Wan, S. Bhosale, S. Zhang, S. Vandenheide, S. Batra, S. Whitman, S. Sootla, S. Collet, S. Gururangan, S. Borodinsky, T. Herman, T. Fowler, T. Sheasha, T. Georgiou, T. Scialom, T. Speckbacher, T. Mihaylov, T. Xiao, U. Karn, V. Goswami, V. Gupta, V. Ramanathan, V. Kerkez, V. Gonguet, V. Do, V. Vogeti, V. Albiero, V. Petrovic, W. Chu, W. Xiong, W. Fu, W. Meers, X. Martinet, X. Wang, X. Wang, X. E. Tan, X. Xia, X. Xie, X. Jia, X. Wang, Y. Goldschlag, Y. Gaur, Y. Babaei, Y. Wen, Y. Song, Y. Zhang, Y. Li,

References II

Y. Mao, Z. D. Coudert, Z. Yan, Z. Chen, Z. Papakipos, A. Singh, A. Srivastava, A. Jain, A. Kelsey, A. Shajnfeld, A. Gangidi, A. Victoria, A. Goldstand, A. Menon, A. Sharma, A. Boesenberg, A. Baevski, A. Feinstein, A. Kallet, A. Sangani, A. Teo, A. Yunus, A. Lupu, A. Alvarado, A. Caples, A. Gu, A. Ho, A. Poulton, A. Ryan, A. Ramchandani, A. Dong, A. Franco, A. Goyal, A. Saraf, A. Chowdhury, A. Gabriel, A. Bharambe, A. Eisenman, A. Yazdan, B. James, B. Maurer, B. Leonhardi, B. Huang, B. Loyd, B. D. Paola, B. Paranjape, B. Liu, B. Wu, B. Ni, B. Hancock, B. Wasti, B. Spence, B. Stojkovic, B. Gamido, B. Montalvo, C. Parker, C. Burton, C. Mejia, C. Liu, C. Wang, C. Kim, C. Zhou, C. Hu, C.-H. Chu, C. Cai, C. Tindal, C. Feichtenhofer, C. Gao, D. Civin, D. Beaty, D. Kreymer, D. Li, D. Adkins, D. Xu, D. Testuggine, D. David, D. Parikh, D. Liskovich, D. Foss, D. Wang, D. Le, D. Holland, E. Dowling, E. Jamil, E. Montgomery, E. Presani, E. Hahn, E. Wood, E.-T. Le, E. Brinkman, E. Arcaute, E. Dunbar, E. Smothers, F. Sun, F. Kreuk, F. Tian, F. Kokkinos, F. Ozgenel, F. Caggioni, F. Kanayet, F. Seide, G. M. Florez, G. Schwarz, G. Badeer, G. Swee, G. Halpern, G. Herman, G. Sizov, Guangyi, Zhang, G. Lakshminarayanan, H. Inan, H. Shojanazeri, H. Zou, H. Wang, H. Zha, H. Habeeb, H. Rudolph, H. Suk, H. Aspegren, H. Goldman, H. Zhan, I. Damlaj, I. Molybog, I. Tufanov, I. Leontiadis, I.-E. Veliche, I. Gat, J. Weissman, J. Geboski, J. Kohli, J. Lam, J. Asher, J.-B. Gaya, J. Marcus, J. Tang, J. Chan, J. Zhen, J. Reizenstein, J. Teboul, J. Zhong, J. Jin, J. Yang, J. Cummings, J. Carvill, J. Shepard, J. McPhie, J. Torres, J. Ginsburg, J. Wang, K. Wu, K. H. U, K. Saxena, K. Khandelwal, K. Zand, K. Matosich, K. Veeraraghavan, K. Michelena, K. Li, K. Jagadeesh, K. Huang, K. Chawla, K. Huang, L. Chen, L. Garg, L. A. L. Silva, L. Bell, L. Zhang, L. Guo, L. Yu, L. Moshkovich, L. Wehrstedt, M. Khabsa, M. Avalani, M. Bhatt, M. Mankus, M. Hasson, M. Lennie, M. Reso, M. Groshev, M. Naumov, M. Lathi, M. Keneally, M. Liu, M. L. Seltzer, M. Valko, M. Restrepo, M. Patel, M. Vyatskov, M. Samvelyan, M. Clark, M. Macey, M. Wang, M. J. Hermoso, M. Metanat, M. Rastegari, M. Bansal, N. Santhanam, N. Parks, N. White, N. Bawa, N. Singhal, N. Egebo, N. Usunier, N. Mehta, N. P. Laptev, N. Dong, N. Cheng, O. Chernoguz, O. Hart, O. Salpekar, O. Kalinli, P. Kent, P. Parekh, P. Saab, P. Balaji, P. Rittner, P. Bontrager, P. Roux, P. Dollar, P. Zvyagina, P. Ratanchandani, P. Yuvraj, Q. Liang, R. Alao, R. Rodriguez, R. Ayub, R. Murthy, R. Nayani, R. Mitra, R. Parthasarathy, R. Li, R. Hogan, R. Battey, R. Wang, R. Howes, R. Rinott, S. Mehta, S. Siby, S. J. Bondu, S. Datta, S. Chugh, S. Hunt, S. Dhillon, S. Sidorov, S. Pan, S. Mahajan, S. Verma, S. Yamamoto, S. Ramaswamy, S. Lindsay, S. Lindsay, S. Feng, S. Lin, S. C. Zha, S. Patil, S. Shankar, S. Zhang, S. Zhang, S. Wang, S. Agarwal, S. Sajuyigbe, S. Chintala, S. Max, S. Chen, S. Kehoe, S. Satterfield, S. Govindaprasad, S. Gupta, S. Deng, S. Cho, S. Virk, S. Subramanian, S. Choudhury, S. Goldman, T. Remez, T. Glaser, T. Best, T. Koehler, T. Robinson, T. Li, T. Zhang, T. Matthews, T. Chou, T. Shaked, V. Vontimitta,

References III

V. Ajayi, V. Montanez, V. Mohan, V. S. Kumar, V. Mangla, V. Ionescu, V. Poenaru, V. T. Mihailescu, V. Ivanov, W. Li, W. Wang, W. Jiang, W. Bouaziz, W. Constable, X. Tang, X. Wu, X. Wang, X. Wu, X. Gao, Y. Kleinman, Y. Chen, Y. Hu, Y. Jia, Y. Qi, Y. Li, Y. Zhang, Y. Zhang, Y. Adi, Y. Nam, Yu, Wang, Y. Zhao, Y. Hao, Y. Qian, Y. Li, Y. He, Z. Rait, Z. DeVito, Z. Rosnbrick, Z. Wen, Z. Yang, Z. Zhao, and Z. Ma, "The Llama 3 Herd of Models," Nov. 2024.

- [5] OpenAI, A. Jaech, A. Kalai, A. Lerer, A. Richardson, A. El-Kishky, A. Low, A. Helyar, A. Madry, A. Beutel, A. Carney, A. Iftimie, A. Karpenko, A. T. Passos, A. Neitz, A. Prokofiev, A. Wei, A. Tam, A. Bennett, A. Kumar, A. Saraiva, A. Vallone, A. Duberstein, A. Kondrich, A. Mishchenko, A. Applebaum, A. Jiang, A. Nair, B. Zoph, B. Ghorbani, B. Zhang, B. Rossen, B. Sokolowsky, B. Barak, B. McGrew, B. Minaiev, B. Hao, B. Baker, B. Houghton, B. McKinzie, B. Eastman, C. Lugaresi, C. Bassin, C. Hudson, C. M. Li, C. de Bourcy, C. Voss, C. Shen, C. Zhang, C. Koch, C. Orsinger, C. Hesse, C. Fischer, C. Chan, D. Roberts, D. Kappler, D. Levy, D. Selsam, D. Dohan, D. Farhi, D. Mely, D. Robinson, D. Tsipras, D. Li, D. Oprica, E. Freeman, E. Zhang, E. Wong, E. Proehl, E. Cheung, E. Mitchell, E. Wallace, E. Ritter, E. Mays, F. Wang, F. P. Such, F. Raso, F. Leoni, F. Tsimpourlas, F. Song, F. von Lohmann, F. Sulit, G. Salmon, G. Parascandolo, G. Chabot, G. Zhao, G. Brockman, G. Leclerc, H. Salman, H. Bao, H. Sheng, H. Andrin, H. Bagherinezhad, H. Ren, H. Lightman, H. W. Chung, I. Kivlichan, I. O'Connell, I. Osband, I. C. Gilaberte, I. Akkaya, I. Kostrikov, I. Sutskever, I. Kofman, J. Pachocki, J. Lennon, J. Wei, J. Harb, J. Twore, J. Feng, J. Yu, J. Weng, J. Tang, J. Yu, J. Q. Candela, J. Palermo, J. Parish, J. Heidecke, J. Hallman, J. Rizzo, J. Gordon, J. Uesato, J. Ward, J. Huizinga, J. Wang, K. Chen, K. Xiao, K. Singhal, K. Nguyen, K. Cobbe, K. Shi, K. Wood, K. Rimbach, K. Gu-Lemberg, K. Liu, K. Lu, K. Stone, K. Yu, L. Ahmad, L. Yang, L. Liu, L. Maksin, L. Ho, L. Fedus, L. Weng, L. Li, L. McCallum, L. Held, L. Kuhn, L. Kondraciuk, L. Kaiser, L. Metz, M. Boyd, M. Trebacz, M. Joglekar, M. Chen, M. Tintor, M. Meyer, M. Jones, M. Kaufer, M. Schwarzer, M. Shah, M. Yatbaz, M. Y. Guan, M. Xu, M. Yan, M. Glaese, M. Chen, M. Lampe, M. Malek, M. Wang, M. Fradin, M. McClay, M. Pavlov, M. Wang, M. Wang, M. Murati, M. Bavarian, M. Rohaninejad, N. McAleese, N. Chowdhury, N. Chowdhury, N. Ryder, N. Tezak, N. Brown, O. Nachum, O. Boiko, O. Murk, O. Watkins, P. Chao, P. Ashbourne, P. Izmailov, P. Zhokhov, R. Dias, R. Arora, R. Lin, R. G. Lopes, R. Gaon, R. Miyara, R. Leike, R. Hwang, R. Garg, R. Brown, R. James, R. Shu, R. Cheu, R. Greene, S. Jain, S. Altman, S. Toizer, S. Toyer, S. Miserendino, S. Agarwal, S. Hernandez, S. Baker, S. McKinney, S. Yan, S. Zhao, S. Hu, S. Santurkar, S. R. Chaudhuri, S. Zhang, S. Fu, S. Papay, S. Lin, S. Balaji, S. Sanjeev, S. Sidor, T. Broda, A. Clark, T. Wang, T. Gordon, T. Sanders, T. Patwardhan, T. Sottiaux, T. Degry, T. Dimson, T. Zheng, T. Garipov, T. Stasi, T. Bansal, T. Creech, T. Peterson, T. Eloundou, V. Qi,

References IV

- V. Kosaraju, V. Monaco, V. Pong, V. Fomenko, W. Zheng, W. Zhou, W. Zhan, W. McCabe, W. Zaremba, Y. Dubois, Y. Lu, Y. Chen, Y. Cha, Y. Bai, Y. He, Y. Zhang, Y. Wang, Z. Shao, and Z. Li, "OpenAI o1 System Card," Apr. 2026.
- [6] A. Jake, "Uber burned through its entire 2026 AI budget in four months. Now its COO is questioning whether it's worth it."
- [7] Q. Cheng, P. Lin, and Y. Zhao, "Does generative AI facilitate investor Trading? Early evidence from ChatGPT outages," *Journal of Accounting and Economics*, vol. 80, p. 101821, Nov. 2025.
- [8] Z. Wang, J. Wohlwend, and T. Lei, "Structured Pruning of Large Language Models," in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, (Online), pp. 6151–6162, Association for Computational Linguistics, 2020.
- [9] G. Hinton, O. Vinyals, and J. Dean, "Distilling the Knowledge in a Neural Network," Mar. 2015.
- [10] J. Wu, Y. Wang, Z. Wu, Z. Wang, A. Veeraraghavan, and Y. Lin, "Deep $\$k\$$ -Means: Re-Training and Parameter Sharing with Harder Cluster Assignments for Compressing Deep Convolutions," June 2018.
- [11] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference," Dec. 2017.
- [12] C. E. Shannon, "A mathematical theory of communication," *The Bell System Technical Journal*, vol. 27, no. 3, pp. 379–423, 1948.
- [13] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, C. Zheng, D. Liu, F. Zhou, F. Huang, F. Hu, H. Ge, H. Wei, H. Lin, J. Tang, J. Yang, J. Tu, J. Zhang, J. Yang, J. Yang, J. Zhou, J. Zhou, J. Lin, K. Dang, K. Bao, K. Yang, L. Yu, L. Deng, M. Li, M. Xue, M. Li, P. Zhang, P. Wang, Q. Zhu, R. Men, R. Gao, S. Liu, S. Luo, T. Li, T. Tang, W. Yin, X. Ren, X. Wang, X. Zhang, X. Ren, Y. Fan, Y. Su, Y. Zhang, Y. Zhang, Y. Wan, Y. Liu, Z. Wang, Z. Cui, Z. Zhang, Z. Zhou, and Z. Qiu, "Qwen3 Technical Report," May 2025.
- [14] Qwen Team, "Qwen3.5: Accelerating productivity with native multimodal agents," Feb. 2026.

References V

- [15] Gemma Team, S. E. Abd, V. Aggarwal, R. Algayres, A. Andreev, O. Bachem, I. Ballantyne, C. Brick, V. Cărbune, M. Casbon, M. Chaturvedi, A. Chawla, V. Cotruta, A. Coucke, P. Culliton, R. Dadashi, L. Dixon, M. Elhawaty, U. Evci, C. Farabet, J. Ferret, F. Galgani, S. Girgin, J.-B. Grill, M. Grootendorst, J. Guo, C. Hardin, Y. He, S. M. Hernandez, O. Homburger, L. Hussenot, J. Ji, A. Joulin, A. Kamath, P. Kassraie, O. Lacombe, P. Lahoti, G. Liu, G. Martins, L. Martins, T. Matejovicova, R. Merhej, N. Momchev, S. Mondal, R. Mullins, S. R. Panyam, S. Pathak, S. Perrin, A. S. Pinto, E. Pot, A. Pouget, A. Ramé, S. Ramos, D. Reid, D. Rim, M. Rivière, K. Roth, L. Rouillard, O. Sanseviero, P. G. Sessa, S. Settle, D. Sinopalnikov, S. Smoot, P. Stanczyk, A. Steiner, L. Stewart, I. Tolstikhin, M. Tschannen, A. Tsitsulin, N. Vieillard, R. Wu, P. Xu, H. Yang, E. Yvinec, B. Zhang, L. Zhang, J. Zou, N. Aagnes, A. Abdelhamed, J. Adamek, S. Agrawal, S. Agrawal, I. Alabdulmohsin, J. B. Alayrac, U. Alon, C. Amarnath, A. Anand, C. Anastasiou, S. Ariafar, F.-X. Aubet, K. Axiotis, F. Barbero, J. Barral, A. Bendebury, U. Bergmann, S. Bileschi, K. Black, M. Blondel, S. Borgeaud, A. Bražińskas, R. Burnell, R. Busa-Fekete, M. Cai, D. Calandriello, G. Cameron, C. Caucheteux, R. Chaabouni, G. Chadha, J. Chan, B. J. Chen, J. Chen, L. Chen, X. Chen, D. Cheng, T.-h. Chien, N. Chinaev, Y. Chou, Z. Chu, B. Coleman, P. Consul, S. Conway-Rahman, S. Crowell, D. Cutler, V. Dani, S. Daruki, A. Das, D. Deutsch, N. Dikkala, L. Ding, Q. Ding, S. Dodhia, K. Donhauser, T. Doshi, A. Dragan, A. Druinsky, S. Dua, Z. Egyed, D. Eisenbud, D. Eppens, C. Fan, B. Fatemi, Y. Fathullah, V. Feinberg, M. Ferev, S. Flennerhag, T. Fujimoto, J. G. Oliveira, I. Galatzer-Levy, J. Gante, S. Geisler, S. Ghosal, A. M. Girgis, T. von Glehn, A. Go, A. Gokhale, A. Grills, Y. Gu, M. Gupta, P. Gupta, G. Guruganesh, R. Hadsell, H. Harkous, J. Harlalka, D. Hassabis, A. Hauth, J. Heyward, A. Hosseini, C.-Y. Hsia, I.-H. Hsu, X. Huang, Y. Huang, K. Hui, A. Hutter, T. I, F. Iliopoulos, A. Jain, G. Jawahar, Z. Ji, Q. Jin, M. Johnson, K. Joshi, A. Kandoor, W.-C. Kang, K. Kavukcuoglu, M. Kazemi, K. Kenealy, A. Khalifa, P. Kirk, I. Korotkov, S. Kothawade, V. Kovalev, N. Kovelamudi, A. Kraft, R. Kumar, V. Kumar, H. Kuppam, J. Lannin, C.-Y. Lee, S. Lee, D. Lepikhin, A. Levkovitch, D. Li, Q. Li, V. Liévin, E. Lin, Z. Lin, C. Liu, T. Liu, T. Liu, X. Liu, I. Lobov, M. Lunayach, M. Ma, G. Madan, A. Maksai, E. Malmi, M. Matuszak, D. McDuff, G. Menghani, M. Miłkoła, D. Mirylenka, K. Misiunas, V. Misra, A. Mitran, K. Mohamed, M. Mukha, E. Noland, J. O'Donnell, B. O'Donoghue, K. Olszewska, B. Orlando, W. Pan, R. Panigrahy, U. Parekh, N. Perez-Nieves, C. Park, E. Paskie, L. Peng, B. Petrini, S. Petrov, J. Pfeiffer, B. Piot, M. Plomecka, S. Poder, O. Ponce, A. Pramanik, D. Racz, A. Rajan, M. Ramanovich, A. Rao, M. Ritter, V. Rodrigues, E. Rosen, M. Rybiński, N. Sachdeva, M. E. Sander, R. Sathyanarayana, S. Savla, S. Schmidgall, T. Schuster, G. Scrivener, B. Seguin, A. Sellergren, A. Severyn, I. Shafran, D. Shah, B. Shahriari, Y. Shangguan, A. Shenoy, P. Shenoy, R. Shivanna, P. Sho, L. Spangher, W. Stokowiec, T. Strother, Y. Su, Y. Sun,

References VI

- M. Sundararajan, A. Tacchetti, M. H. Taege, P. Tafti, J. Tarbouriech, C. Tekur, S. Thakoor, R. Thapa, M. Traverse, L. Treven, T. Tu, C. T. Tung, Ç. Ünlü, P. Veličković, M. P. Venkat, S. G. Venkatesh, V. Venkiteswaran, F. Visin, A. Vitvitskyi, K. Vodrahalli, W. Wang, X. Wang, T. Warkentin, J. Wassenberg, J. Wieting, C. Wu, L. Xiao, H. Xu, Y. Xu, F. Xue, A. Yadav, J. Yan, A. Yang, L. Yang, M.-H. Yang, Z. Ying, J. H. Yoo, M. Zadimoghaddam, S. Zafar, F. Zhang, J. Zhang, J. Zhang, X. Zhang, C. Zhao, D. Zhou, and C. Zou, "Gemma 4 Technical Report," July 2026.
- [16] Z.ai inc., "GLM-4.7: Advancing the Coding Capability," Dec. 2025.
- [17] R. Eldan and Y. Li, "TinyStories: How Small Can Language Models Be and Still Speak Coherent English?," May 2023.
- [18] S. Black, L. Gao, P. Wang, C. Leahy, and S. Biderman, "GPT-Neo: Large scale autoregressive language modeling with meshtensorflow." Zenodo, Oct. 2021.
- [19] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, "Language models are unsupervised multitask learners," 2019.
- [20] T. Kudo and J. Richardson, "SentencePiece: A simple and language independent subword tokenizer and detokenizer for Neural Text Processing," Aug. 2018.
- [21] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, A. Rodriguez, A. Joulin, E. Grave, and G. Lample, "LLaMA: Open and Efficient Foundation Language Models," Feb. 2023.
- [22] F. Jelinek, R. L. Mercer, L. R. Bahl, and J. K. Baker, "Perplexity—a measure of the difficulty of speech recognition tasks," *The Journal of the Acoustical Society of America*, vol. 62, pp. S63–S63, Dec. 1977.
- [23] M. Nagel, R. A. Amjad, M. van Baalen, C. Louizos, and T. Blankevoort, "Up or Down? Adaptive Rounding for Post-Training Quantization," June 2020.
- [24] J. M. Cohen, S. Kaur, Y. Li, J. Z. Kolter, and A. Talwalkar, "Gradient Descent on Neural Networks Typically Occurs at the Edge of Stability," Nov. 2022.

References VII

- [25] Kimi Team, Y. Bai, Y. Bao, Y. Charles, C. Chen, G. Chen, H. Chen, H. Chen, J. Chen, N. Chen, R. Chen, Y. Chen, Y. Chen, Y. Chen, Z. Chen, J. Cui, H. Ding, M. Dong, A. Du, C. Du, D. Du, Y. Du, Y. Fan, Y. Feng, K. Fu, B. Gao, C. Gao, H. Gao, P. Gao, T. Gao, Y. Ge, S. Geng, Q. Gu, X. Gu, L. Guan, H. Guo, J. Guo, X. Hao, T. He, W. He, W. He, Y. He, C. Hong, H. Hu, Y. Hu, Z. Hu, W. Huang, Z. Huang, Z. Huang, T. Jiang, Z. Jiang, X. Jin, Y. Kang, G. Lai, C. Li, F. Li, H. Li, M. Li, W. Li, Y. Li, Y. Li, Y. Li, Z. Li, Z. Li, H. Lin, X. Lin, Z. Lin, C. Liu, C. Liu, H. Liu, J. Liu, J. Liu, L. Liu, S. Liu, T. Y. Liu, T. Liu, W. Liu, Y. Liu, Y. Liu, Y. Liu, Z. Liu, E. Lu, H. Lu, L. Lu, Y. Luo, S. Ma, X. Ma, Y. Ma, S. Mao, J. Mei, X. Men, Y. Miao, S. Pan, Y. Peng, R. Qin, Z. Qin, B. Qu, Z. Shang, L. Shi, S. Shi, F. Song, J. Su, Z. Su, L. Sui, X. Sun, F. Sung, Y. Tai, H. Tang, J. Tao, Q. Teng, C. Tian, C. Wang, D. Wang, F. Wang, H. Wang, H. Wang, J. Wang, J. Wang, J. Wang, S. Wang, S. Wang, S. Wang, X. Wang, Y. Wang, Y. Wang, Y. Wang, Y. Wang, Y. Wang, Z. Wang, Z. Wang, Z. Wang, Z. Wang, C. Wei, Q. Wei, H. Wu, W. Wu, X. Wu, Y. Wu, C. Xiao, J. Xie, X. Xie, W. Xiong, B. Xu, J. Xu, L. H. Xu, L. Xu, S. Xu, W. Xu, X. Xu, Y. Xu, Z. Xu, J. Xu, J. Xu, J. Yan, Y. Yan, H. Yang, X. Yang, Y. Yang, Y. Yang, Z. Yang, Z. Yang, Z. Yang, H. Yao, X. Yao, W. Ye, Z. Ye, B. Yin, L. Yu, E. Yuan, H. Yuan, M. Yuan, S. Yuan, H. Zhan, D. Zhang, H. Zhang, W. Zhang, X. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Z. Zhang, H. Zhao, Y. Zhao, Z. Zhao, H. Zheng, S. Zheng, L. Zhong, J. Zhou, X. Zhou, Z. Zhou, J. Zhu, Z. Zhu, W. Zhuang, and X. Zu, "Kimi K2: Open Agentic Intelligence," Feb. 2026.
- [26] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, "Scikit-learn: Machine learning in python," *Journal of Machine Learning Research*, vol. 12, no. 85, pp. 2825–2830, 2011.
- [27] E. Borenstein, E. Sharon, and S. Ullman, "Combining top-down and bottom-up segmentation," vol. 2004, pp. 46–46, July 2004.
- [28] D. S. Kermamy, M. Goldbaum, W. Cai, C. C. Valentim, H. Liang, S. L. Baxter, A. McKeown, G. Yang, X. Wu, F. Yan, J. Dong, M. K. Prasadha, J. Pei, M. Y. Ting, J. Zhu, C. Li, S. Hewett, J. Dong, I. Ziyar, A. Shi, R. Zhang, L. Zheng, R. Hou, W. Shi, X. Fu, Y. Duan, V. A. Huu, C. Wen, E. D. Zhang, C. L. Zhang, O. Li, X. Wang, M. A. Singer, X. Sun, J. Xu, A. Tafreshi, M. A. Lewis, H. Xia, and K. Zhang, "Identifying Medical Diagnoses and Treatable Diseases by Image-Based Deep Learning," *Cell*, vol. 172, pp. 1122–1131.e9, Feb. 2018.
- [29] I. Loshchilov and F. Hutter, "Decoupled Weight Decay Regularization," Jan. 2019.
- [30] I. Loshchilov and F. Hutter, "SGDR: Stochastic Gradient Descent with Warm Restarts," May 2017.