

Gradient-Aware Adaptive Rounding for Post-Training Quantization of Deep Neural Networks

48246347 Piye Liu

Supervisor: Prof. Hiroshi Nakamura

The recent exponential scaling of Large Language Models (LLMs) has resulted in massive carbon emissions and significant challenges in deployment, forcing users to rely on third-party cloud providers, thereby raising privacy concerns. This highlights the necessity of model compression, and post-training quantization (PTQ) is one method that enables efficient on-device models. AdaRound, an adaptive rounding method that outperforms round-to-nearest rounding during quantization, minimizes the task loss between the original weights and the quantized weights on unlabelled calibration data. We find that: 1) AdaRound leaks some information about its calibration data into the weight perturbation matrix, and we demonstrate this by proposing a novel inversion attack that targets the null space of this matrix to recover coarse representations of the original data. 2) The theoretical foundation of AdaRound asserts that the gradient of the loss function is zero for pretrained converged models, an assumption which is refuted by recent literature. To address this discrepancy, we attempted to reconcile this by incorporating local gradient signals into the AdaRound optimization process, but obtained neutral results on empirical evaluations. Therefore, this thesis highlights a frequently overlooked dimension in PTQ (privacy), and proposes a gradient-aware modification of AdaRound to better align the method with modern literature.

Outline

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Exponential Scaling of LLMs

Recently, the exponential scaling of LLM sizes has inadvertently resulted in models that are extremely challenging to deploy, thus causing users to rely on third-party providers for LLM usage.




Figure 1: Observed exponential scaling of LLM sizes over time [1] (Note: The GPT-4 figure is an estimate)

Why Compress LLMs?

- ▶ Environmental Effects: GPT-3's training had an estimated CO₂e of ~ 552.1 tons [2], GPT-4's training had an estimated CO₂e of ~ 6912 tons [3], and Llama 3.3 [4]'s training had an estimated CO₂e of $\sim 11,390$ tons.
- ▶ Cost: For example, OpenAI o1 [5] in “pro mode” costs \$150 per 1 million tokens input, and \$600 per 1 million tokens output. Uber's yearly AI budget was exhausted in 4 months, while Microsoft began cancelling Claude Code licenses just 6 months after acquiring them [6].
- ▶ Reliability. Single point of failure, and an outage can have severe effects on mission-critical tasks. For example, it is observed that whenever ChatGPT has an outage, stock trading volume goes down [7].

Benefits of On-Device Inference

- ▶ Privacy: Data does not need to leave the device
- ▶ Robustness: Can function even without internet access
- ▶ Latency: Doesn't have to rely on third-party scheduling, can prioritize urgent tasks
- ▶ Cost: Reduces API fees, network costs, and even environmental emissions

Model Compression Methods

Methods typically used for the compression of deep neural networks:

- ▶ Weight Pruning: Setting some weights to zeroes, resulting in sparse weight matrices [8].
- ▶ Knowledge Distillation: Smaller model learns to match output logits of larger model [9].
- ▶ Weight Palettization: Clusters weights (e.g. via k -means), then stores every weight as the centroid of the cluster they belong to [10].
- ▶ Quantization: **We focus on this method for our thesis.** Stores weights in a data type of less bit width (e.g. int8), compared to the original data type (e.g. float32) [11].

Contributions

Our contributions in this work are twofold:

- ▶ First, we demonstrate a novel inversion method against AdaRound to recover coarse outlines that represent its calibration data, motivating the need to consider privacy in data-dependent quantization method.
- ▶ Second, we propose a gradient-aware adaptive rounding method based on AdaRound, and apply the proposed method on two language models.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Compressibility of Deep Neural Networks

Source coding theorem [12]:

- ▶ n i.i.d. (independent and identically distributed) random variables, each with entropy $H(X)$, can be compressed to $nH(X)$ bits on average, with negligible information loss, as $n \rightarrow \infty$

In this case, $H(X)$, the entropy of a random variable, is the expected amount of information needed to describe it, and can be computed by:

$$H(X) := - \sum_{x \in X} p(x) \log_2 p(x). \quad (1)$$

Note that the log base determines the units: $\log_2$ yields bits, $\log_{10}$ yields dits, and $\ln$ yields nats.

Measuring Compressibility of LLMs

Total information stored in LLM weights: $\sim n \cdot \hat{H}(W)$ bits;
Maximum possible total entropy of the model: $n \cdot b$ bits. It therefore follows that the capacity utilization by a model is:

$$U(W) := \frac{\hat{H}(W)}{b}. \quad (2)$$

We use this method to quantify the capacity utilization of several modern LLMs:

Table 1: Entropy Estimate of Language Models (bfloat16)

Model	$\hat{H}(W)$ in bits	$U(W)$
Qwen 3 4B [13]	10.58	$10.58/16 \approx 66\%$
Qwen 3.5 4B [14]	10.56	$10.56/16 \approx 66\%$
Gemma 4 12B Instruct [15]	10.69	$10.69/16 \approx 67\%$
GLM 4.7 Flash [16] (30B)	10.49	$10.49/16 \approx 66\%$

Small Language Models

TinyStories [17] is a work that proposed a synthetic low-vocabulary dataset and several small language models, as low as 1 million parameters, that produce coherent English.

One limitation we observed in this work is their usage of the default GPT-Neo [18] tokenizer (which is identical to the pretrained GPT-2 [19] tokenizer), restricted to the top 10,000 tokens.

As a preliminary experiment, we train two small language models on TinyStories v1 and TinyStories v2, respectively. We use a SentencePiece Byte-Pair Encoding (BPE) tokenizer [20] with a vocabulary size of 1024, and we used a LLaMa-based [21] model with roughly half the parameters of the smallest reference model.

Evaluation of Our LMs

We use perplexity [22] as a quantitative metric to evaluate a language model's accuracy. For a model θ , a text x of length N , it is defined as the following.

$$\text{Perplexity}(\theta, x) = \left(\prod_{i=1}^N \frac{1}{p_{\theta}(x_i \mid x_1, x_2, \dots, x_{i-1})} \right)^{\frac{1}{N}} \quad (3)$$

In other words, it is the geometric mean of the inverse probabilities of a given text.

Table 2: Evaluation of Our Language Models

	Reference	Ours (TinyStories v1)	Ours (TinyStories v2)
Architecture	GPT-Neo	LLaMa	LLaMa
Parameter Count	1.17M	0.57M	0.46M
Training Data	TinyStories v1		TinyStories v2
Tokenizer (Vocab Size)	GPT-Neo	SentencePiece	SentencePiece
Vocab Size	10000	1024	1024
Perplexity $\downarrow$ (TinyStories v1)	19.3919	11.9006	—
Perplexity $\downarrow$ (TinyStories v2)	—	—	8.0994

Baseline Round-to-Nearest Quantization

Consider the symmetric uniform quantization method to int8:

- ▶ Quantization:

- ▶ $\text{scale} = \frac{\max(|W|)}{2^7 - 1}$

- ▶ $W_{\text{quantized}} = \text{clamp}(\lfloor \frac{W}{\text{scale}} \rceil)$

- ▶ Reconstruction (dequantization):

- ▶ $W_{\text{reconstructed}} = \text{scale} \cdot W_{\text{quantized}}$

- ▶ Note that $\text{clamp}()$ here refers to clamping a number between $[-2^7, 2^7 - 1]$ to ensure its within the range of the data type.

We can see that the model weights are perturbed in the direction of $W_{\text{reconstructed}} - W$ after quantization and dequantization. Let us use our TinyStories v2 model to investigate the effects of perturbation direction on the model's accuracy.

Loss Landscape: Straight Line

First, we will sample models along a simple line:

$$\theta' = \theta + \epsilon, \epsilon \in \{-0.0020 \dots 0.0015\} \quad (4)$$

Note that the blue line in the figure represents $\epsilon \in \{-0.0015 \dots 0.0015\}$.

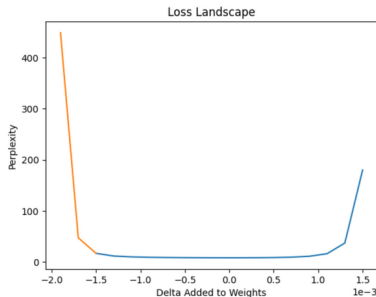


Figure 2: Models Along a Line

Loss Landscape: Moving Randomly

Next, we sample models stochastically from the following distribution. We sample 5 models for each σ .

$$\theta' \sim \mathcal{N}(\theta, \sigma^2 \mathbf{I}), \sigma^2 \in \{0.0005 \dots 0.0040\} \quad (5)$$

Note that we cap the perplexity at 100 in Figure 3 for better visualization, because some perplexities are extremely large.

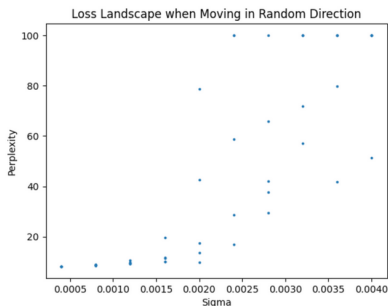
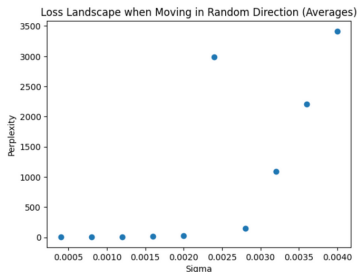


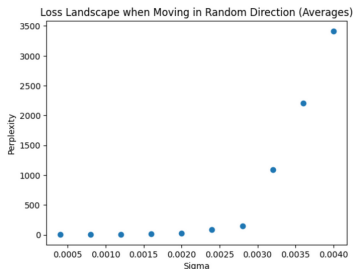
Figure 3: Models from Gaussian Distribution

Loss Landscape: Moving Randomly (Continued)

In Figure 4a, we plot the average perplexities of each σ . We notice an outlier at $\sigma = 0.0024$, and removing it yields Figure 4b.



(a) Averages



(b) Averages (Outlier Removed)

Figure 4: Loss Landscape of Models from Gaussian Distribution

Loss Landscape: Moving Towards Another Model

We pretrain another model with the same configuration as the base model, and move from the base model to the other model. The distance between these models, $\|\theta - \theta'\|_2$, is 75.6, which means they're quite far apart, and the magnitude of perturbation is much greater than the previous experiments.

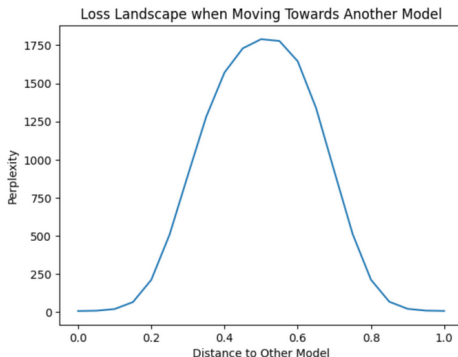


Figure 5: Moving Towards Another Model

Loss Landscape: Observations

- ▶ When moving along a line, we notice that there is a “basin” of low loss. Moving towards the negative direction caused less degradation than moving towards the positive direction.
- ▶ There was a huge outlier when moving randomly.
- ▶ Moving towards another model resulted in less degradation despite the massive distance.

We therefore find that: **When weights of a model is perturbed, the loss does not degrade uniformly.**

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

AdaRound: Objective

AdaRound [23] is a novel rounding method, which aims to minimize the difference in task loss of the quantized model and the unquantized model. The task loss of layer j of the unquantized model is $\mathcal{L}(x, y, w_j)$, and the task loss of the same layer of the quantized model $\mathcal{L}(x, y, w_j + \Delta w_j)$. Therefore, the objective is:

$$\arg \min_{\Delta w_j} (\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j)). \quad (6)$$

Using second order Taylor expansion:

$$\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j) \approx \nabla \mathcal{L}(x, y, w_j)^T \Delta w_j + \frac{1}{2} \Delta w_j^T \mathbf{H} \Delta w_j \quad (7)$$

AdaRound: Approximations

The authors assert that $\nabla \mathcal{L}(x, y, w_j)$ is negligible, which transforms the objective into just:

$$\mathcal{L}(x, y, w_j + \Delta w_j) - \mathcal{L}(x, y, w_j) \approx \frac{1}{2} \Delta w_j^T \mathbf{H} \Delta w_j \quad (8)$$

Finally, the authors formulate the optimization of weigh rounding for layer j as the following Quadratic Unconstrained Binary Optimization (QUBO) problem.

$$\arg \min_{\Delta w_j} \Delta w_j^T \mathbf{H} \Delta w_j \quad (9)$$

The above problem is NP-hard, so AdaRound uses the local mean squared error alongside a regularizer as a proxy.

$$\arg \min_{\mathbf{V}} \left(\|\mathbf{W}_X - \hat{\mathbf{W}}_X\|_F^2 + f_{\text{reg}}(\mathbf{V}) \right) \quad (10)$$

AdaRound: Algorithm

Algorithm 1 AdaRound

Require: $\mathbf{W}$, X (calibration data), t (number of iterations), b (target bit width), β (regularization parameter)

Ensure: Quantized model $\mathbf{W}_Q$, scales

- 1: $x_{in} \leftarrow \text{CACHEACTIVATIONS}(\mathbf{W}, X)$ ▷ Forward pass on calibration data to cache input activations
 - 2: **for** w_j in $\mathbf{W}$ **do**
 - 3: $q_{\min} \leftarrow -2^{b-1}$
 - 4: $q_{\max} \leftarrow 2^{b-1} - 1$
 - 5: $\text{scales}_j \leftarrow \frac{\max(w_j)}{q_{\max}}$ ▷ Store scale for output
 - 6: $\mathbf{V}_j \leftarrow h^{-1}(w_j - \lfloor w_j \rfloor)$ ▷ Initialize $\mathbf{V}_j$
 - 7: **for** $i \in \{0 \dots t\}$ **do**
 - 8: $\beta \leftarrow \text{ANNEAL}(\beta)$
 - 9: $\hat{w}_j \leftarrow \text{clamp}\left(\lfloor \frac{w_j}{\text{scales}_j} \rfloor + h(\mathbf{V}_j), q_{\min}, q_{\max}\right)$
 - 10: $\mathcal{L} \leftarrow \|w_j x_{in} - \hat{w}_j x_{in}\|_F^2 + \sum_{i,j} (1 - |2h(\mathbf{V}_{i,j}) - 1|^\beta)$
 - 11: $\mathbf{V} \leftarrow \mathbf{V} - \nabla \mathcal{L}$ ▷ via autograd
 - 12: **end for**
 - 13: $\mathbf{W}_{Q_{i,j}} \leftarrow \begin{cases} \lceil w_{i,j} \rceil, & \text{where } h(\mathbf{V}_{i,j}) \geq 0.5 \\ \lfloor w_{i,j} \rfloor, & \text{otherwise} \end{cases}$ ▷ Finalize quantization
 - 14: **end for**
 - 15: **return** $\mathbf{W}_Q$, scales
-

Limitations of AdaRound

- ▶ First, using calibration data can cause some information to be encoded into the weight perturbation matrix. We empirically verify this in Contribution 1.
- ▶ Second, the assumption in Equation 7 that the gradient of the loss function for pretrained converged models is negligible (for which empirical evidence was not provided) is challenged by recent literature.

AdaRound's Zero-Gradient Assumption

While one might expect gradient descent to follow a smooth, stable path to an optimum, in reality it almost always operate in a highly dynamic, oscillatory regime described as the **Edge of Stability** [24], leading to chaotic loss curves such as:

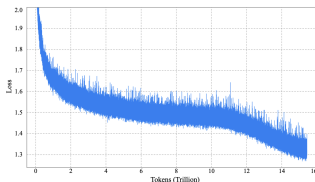


Figure 6: Training Loss Curve of Kimi K2 [25]

The oscillatory optimization also implies that the gradient for pretrained models are rarely zero. Recent literature [24] empirically show that even as the training loss decreases, the gradient norm of a model almost stays constant.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Leakage Channel: The Null Space

- ▶ If we look at the AdaRound objective, the goal is to make $\|Wx - \hat{W}x\|_F^2$ approach zero. If we define the weight perturbation as $\Delta W = W - \hat{W}$, this means that for a model quantized using AdaRound:

$$\Delta Wx \approx 0$$

- ▶ Mathematically, this implies that the calibration data x is approximately in the **null space** of the perturbation matrix. In other words, $x \in \mathcal{N}(\Delta W)$. We can therefore exploit this property to generate samples that resemble the mean of the calibration data.
- ▶ We **can't** reconstruct exact pixel-for-pixel calibration samples. However, we demonstrate that by implementing this as a simple optimization problem, we can reconstruct representations that look like the mean of these calibration samples.

Realizing an Attack

- ▶ If an attacker has access to both the original weights (W) and the quantized weights ($\hat{W}$), they can recover the calibration data by searching for any vector that satisfies this null space property.
- ▶ For example, hypothetically: a company downloads an open-source vision model, quantizes it with AdaRound using confidential internal data (e.g. patient records), then deploys it onto edge devices sold to customers. An attacker can reconstruct representations of the calibration data by extracting the quantized $\hat{W}$ from the edge devices, then performing the inversion on the weight perturbation matrix $\Delta W = W - \hat{W}$.

Proposed Inversion Algorithm

Algorithm 2 AdaRound Inversion Algorithm

Require: $W, \hat{W}, \hat{x}_{\text{init}}$, No. of Iterations K , Learning Rate η

1: $\Delta W \leftarrow W - \hat{W}$

2: $\hat{x} \leftarrow \hat{x}_{\text{init}}$

3: **for** $k = 1$ **to** K **do**

4: $\mathcal{L} \leftarrow \sum (\Delta W \tanh(\hat{x}))^2$ (Forward pass)

5: $\frac{\partial \mathcal{L}}{\partial \hat{x}} \leftarrow \text{Autograd}(\mathcal{L}, \hat{x})$ (Backward pass)

6: $\hat{x} \leftarrow \hat{x} - \eta \frac{\partial \mathcal{L}}{\partial \hat{x}}$ (Update $\hat{x}$)

7: **end for**

8: **return** $\tanh(\hat{x})$

Simple Optimization

Note that the objective can be rewritten as:

$$\|\Delta W \hat{x}\|_2^2 = (\Delta W \hat{x})^T (\Delta W \hat{x}) = \hat{x}^T (\Delta W^T \Delta W) \hat{x}. \quad (11)$$

We can therefore make this simple optimization:

Algorithm 3 Optimized AdaRound Inversion Algorithm

Require: $W, \hat{W}, \hat{x}_{\text{init}},$ No. of Iterations K , Learning Rate η

- 1: $\Delta W \leftarrow W - \hat{W}$
 - 2: $\hat{x} \leftarrow \hat{x}_{\text{init}}$
 - 3: $G \leftarrow \Delta W^T \Delta W$
 - 4: **for** $k = 1$ **to** K **do**
 - 5: $\mathcal{L} \leftarrow \sum (\tanh(\hat{x})^T G \tanh(\hat{x}))$ (Forward pass)
 - 6: $\frac{\partial \mathcal{L}}{\partial \hat{x}} \leftarrow \text{Autograd}(\mathcal{L}, \hat{x})$ (Backward pass)
 - 7: $\hat{x} \leftarrow \hat{x} - \eta \frac{\partial \mathcal{L}}{\partial \hat{x}}$ (Update $\hat{x}$)
 - 8: **end for**
 - 9: **return** $\tanh(\hat{x})$
-

Experiment Setup

- ▶ We attempt an inversion on MLPs trained on the Olivetti faces dataset [26] the Weizmann Horse Database[27], and a Chest X-Ray Images dataset[28] using an MLP classifier quantized by AdaRound.
- ▶ For each model, we performed 500 iterations of AdaRound quantization. We then attempted to invert the process by training 10 random Gaussian samples for 3,000 iterations using the Adam optimizer. We report SSIM, PSNR, and Normalized RMSE before and after the inversion.

Results

Table 3: Metrics from Random Initialization vs. Our Method

	SSIM $\uparrow$	PSNR $\uparrow$	NRMSE $\downarrow$
Olivetti Faces (Before)	0.00	3.50	1.16
Olivetti Faces (After)	0.14 (+0.14)	12.83 (+9.33)	0.40 (-0.76)
Chest X-Ray (Before)	0.00	2.61	1.10
Chest X-Ray (After)	0.07 (+0.07)	10.13 (+7.52)	0.46 (-0.64)
Horse DB (Before)	0.00	2.60	1.11
Horse DB (After)	0.07 (+0.07)	10.15 (+7.55)	0.47 (-0.64)

The increase in PSNR and the decrease in RMSE across all datasets (average over 10 samples) indicate that our method is recovering real structural information, although it is very coarse.

Ablation Study

- ▶ We perform the same inversion method against the same models quantized with round-to-nearest (RTN) rounding.
- ▶ Since RTN does not use calibration data, we compare our method's samples to the mean of the training data instead.

Table 4: Results from Ablation Study

	SSIM $\uparrow$	PSNR $\uparrow$	NRMSE $\downarrow$
Olivetti Faces (Before)	0.00	2.89	1.12
Olivetti Faces (After)	0.00	3.90 (+1.01)	1.00 (-0.12)
Chest X-Ray (Before)	0.00	2.77	1.11
Chest X-Ray (After)	0.00	3.64 (+0.87)	1.00 (-0.11)
Horse DB (Before)	0.00	2.57	1.10
Horse DB (After)	0.00	3.38 (+0.81)	1.00 (-0.10)

- ▶ We notice 0.00 SSIM on all datasets, implying that the generated samples do not look like the calibration data at all.

Discussion & Limitations

- ▶ Even though the SSIM in the main experiment is quite low, we observe empirical evidence of the leakage channel. Furthermore, the SSIM in the ablation study is 0.00 on all datasets, which empirically shows that AdaRound is likely to be the source of the leakage channel.
- ▶ The main limitation is that our method only reconstructs samples that look like the mean of the training data. It does not reconstruct individual samples as-is. Furthermore, the SSIM we obtained is rather low and there is much room for further improvement.
- ▶ Another limitation is that to perform this attack, the weight perturbation matrix is required. Typically this means needing to obtain both the unquantized and quantized models.

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Method

- ▶ To reconcile AdaRound's assumption that the gradient term in Equation 7 can be ignored, and Central Flows' evidence that the gradient is nonzero even for pretrained converged models, we propose a modified version of AdaRound to incorporate awareness of the local gradient.
- ▶ In order to minimize

$$\mathbb{E} \left[g^{(w)^T} \Delta W \right], \quad (12)$$

we can attempt to force ΔW to have the same sign as $-g^{(w)^T}$.

- ▶ We do this by finetuning the model on a small amount of data then taking the EMA of the model during finetuning as W'_{avg} . We add the Binary Cross Entropy (BCE) between the current rounding directions and $\text{sign}(W'_{\text{avg}} - W)$ as a penalty term to AdaRound's optimization objective.

Algorithm for EMA Finetuning

Algorithm 4 Gradient-Aware Adaptive Rounding (Finetuning)

Require: $\mathbf{W}$, X (calibration data), η (learning rate), B (batch size)

Ensure: EMA of models during finetuning $\mathbf{W}'_{\text{avg}}$

- 1: $\mathbf{W}' \leftarrow \mathbf{W}$
 - 2: $\mathbf{W}'_{\text{avg}} \leftarrow \mathbf{W}$
 - 3: $N \leftarrow \frac{X}{B}$ ▷ Number of Iterations
 - 4: $\alpha \leftarrow \frac{2}{N+1}$ ▷ EMA Smoothing Factor
 - 5: **for** each batch x_b in X **do**
 - 6: $\hat{x}_b \leftarrow \mathbf{W}'(x_b)$ ▷ Forward pass
 - 7: $\mathcal{L} \leftarrow \text{CrossEntropyLoss}(\hat{x}_b, x_b)$
 - 8: $\mathbf{W}' \leftarrow \mathbf{W}' - \eta \cdot \nabla \mathcal{L}$ ▷ Backward pass
 - 9: $\mathbf{W}'_{\text{avg}} \leftarrow \alpha \cdot \mathbf{W}'_{\text{avg}} + (1 - \alpha) \cdot \mathbf{W}'$ ▷ Update EMA
 - 10: **end for**
 - 11: **return** $\mathbf{W}'_{\text{avg}}$
-

Proposed Method Objective

Our proposed method is essentially a modification to AdaRound's objective in Equation 10:

$$\arg \min_{\mathbf{V}} \left(\|\mathbf{W}_X - \hat{\mathbf{W}}_X\|_F^2 + f_{\text{reg}}(\mathbf{V}) + \underbrace{\psi \cdot \text{BCE} \left(h(\mathbf{V}_j), \text{sign}(\mathbf{W}' - \mathbf{W}) \right)}_{\text{Proposed addition}} \right) \quad (13)$$

Experiment Setup

We evaluate the proposed method on our TinyStories models with the following hyperparameters and components.

Table 5: Implementation Details

Hyperparameter/Component	Value
Finetuning Data	2% of Training Data
Learning Rate	3e-4
Optimizer	AdamW [29]
Learning Rate Scheduler	CosineAnnealingLR [30]
ψ	100.0

Results: Main Experiment

- ▶ We observe **neutral to negative** results using the proposed method.
- ▶ Note that the two models had different testing data (TinyStories v1 and TinyStories v2, respectively).

Table 6: Perplexity Evaluation of Proposed Method

Type of Quantization	TinyStories V1	TinyStories V2
None (Unquantized Model)	11.9006	8.0994
Round-to-Nearest	11.9219	8.1224
AdaRound	11.9216	8.1177
Proposed Method	11.9216	8.1179

Ablation Study: Method

We report an ablation study by ablating two possible elements of the proposed method: AdaRound and EMA. Ablating EMA means we use final finetuned model weights instead of the EMA of model weights during all iterations. Ablating AdaRound means we steer the rounding using only the sign of the weight perturbation matrix:

Algorithm 5 Gradient-based Steering without AdaRound

Require: $\mathbf{W}$ (unquantized weights), $\mathbf{W}'$ (calibration model), b (target bit width)

Ensure: Quantized model $\mathbf{W}_q$

```
1: function CUSTOMROUND( $\mathbf{W}$ ,  $\mathbf{W}'$ )
2:    $\mathbf{W}_{i,j} \leftarrow \begin{cases} \lceil \mathbf{W}_{i,j} \rceil, & \text{where } \mathbf{W}_{i,j} < \mathbf{W}'_{i,j} \\ \lfloor \mathbf{W}_{i,j} \rfloor, & \text{otherwise} \end{cases}$ 
3:   return  $\mathbf{W}$ 
4: end function
5:  $q_{\min} \leftarrow -2^{b-1}$ 
6:  $q_{\max} \leftarrow 2^{b-1} - 1$ 
7:  $\text{scale} \leftarrow \frac{\max(\mathbf{W})}{q_{\max}}$ 
8:  $\mathbf{W}_q \leftarrow \text{clamp}\left(\text{CUSTOMROUND}\left(\frac{\mathbf{W}}{\text{scale}}\right), q_{\min}, q_{\max}\right)$ 
9: return  $\mathbf{W}_q$ , scale
```

Ablation Study: Results

Table 7: Perplexity Evaluation of Ablation Studies

	AdaRound + BCE	EMA	TinyStories V1	TinyStories V2
Main Experiment	✓	✓	11.9216	8.1179
Ablation 1	✓	×	11.9216	8.1197
Ablation 2	×	✓	11.9166	8.1546
Ablation 3	×	×	11.9050	8.1084

- ▶ We notice that we achieved the best performance by ablating both EMA and AdaRound (outperforming original AdaRound)
- ▶ This ablation study shows that the gradient direction can still potentially be useful despite the neutral to negative results in the main experiment.

Discussion

1. Motivation
2. Preliminaries
3. Literature Review
4. Contribution 1: Inversion Attack on AdaRound
5. Contribution 2: Gradient-Aware Adaptive Rounding
6. Conclusion

Conclusion

- ▶ We first quantified the compressibility of modern LLMs and confirmed that their capacity is underutilized.
- ▶ We analysed the loss landscape of a language model to show the effects of weight perturbation direction on a model's accuracy.
- ▶ We find that AdaRound inadvertently encodes information about its calibration data into its weight perturbation matrix, and verified empirically.
- ▶ We challenged an assertion in AdaRound that the gradient of the loss function for pretrained converged models is negligible. However, we obtained neutral to negative results using our proposed method.

Future Work

In the future, we would like to:

- ▶ Improve the quality of the inversion attack using regularizers (e.g. Total Variation (TV) loss) and priors (e.g. deep inversion priors).
- ▶ Furthermore explore the incorporation of the gradient of the loss function in AdaRound, possibly using other proxies to estimate the direction of the gradient.
- ▶ Replicate our experiments on models of larger scale, so we can determine whether its behavior persists beyond smaller on-device models.

References

- [1] J. Gerstmayr, P. Manzl, and M. Pieber, "Multibody models generated from natural language," *Multibody System Dynamics*, vol. 62, pp. 249–271, Jan. 2024.
- [2] D. Patterson, J. Gonzalez, Q. Le, C. Liang, L.-M. Munguia, D. Rothchild, D. So, M. Texier, and J. Dean, "Carbon Emissions and Large Neural Network Training," Apr. 2021.
- [3] Z. Kausar, S. Latif, R. K. Shahzad, and M. Fatima, "Quantifying the Climate Risk of Generative AI: Region-Aware Carbon Accounting with G-TRACE and the AI Sustainability Pyramid," May 2026.
- [4] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, A. Yang, A. Fan, A. Goyal, A. Hartshorn, A. Yang, A. Mitra, A. Sravankumar, A. Korenev, A. Hinsvark, A. Rao, A. Zhang, A. Rodriguez, A. Gregerson, A. Spataru, B. Roziere, B. Biron, B. Tang, B. Chern, C. Caucheteux, C. Nayak, C. Bi, C. Marra, C. McConnell, C. Keller, C. Touret, C. Wu, C. Wong, C. C. Ferrer, C. Nikolaidis, D. Allonsius, D. Song, D. Pintz, D. Livshits, D. Wyatt, D. Esiobu, D. Choudhary, D. Mahajan, D. Garcia-Olano, D. Perino, D. Hupkes, E. Lakomkin, E. AlBadawy, E. Lobanova, E. Dinan, E. M. Smith, F. Radenovic, F. Guzmán, F. Zhang, G. Synnaeve, G. Lee, G. L. Anderson, G. Thattai, G. Nail, G. Mialon, G. Pang, G. Cucurell, H. Nguyen, H. Korevaar, H. Xu, H. Touvron, I. Zarov, I. A. Ibarra, I. Kloumann, I. Misra, I. Evtimov, J. Zhang, J. Copet, J. Lee, J. Geffert, J. Vranes, J. Park, J. Mahadeokar, J. Shah, J. van der Linde, J. Billock, J. Hong, J. Lee, J. Fu, J. Chi, J. Huang, J. Liu, J. Wang, J. Yu, J. Bitton, J. Spisak, J. Park, J. Rocca, J. Johnstun, J. Saxe, J. Jia, K. V. Alwala, K. Prasad, K. Upasani, K. Plawiak, K. Li, K. Heafield, K. Stone, K. El-Arini, K. Iyer, K. Malik, K. Chiu, K. Bhalla, K. Lakhota, L. Rantala-Young, L. van der Maaten, L. Chen, L. Tan, L. Jenkins, L. Martin, L. Madaan, L. Malo, L. Blecher, L. Landzaat, L. de Oliveira, M. Muzzi, M. Pasupuleti, M. Singh, M. Paluri, M. Kardas, M. Tsimpoukelli, M. Oldham, M. Rita, M. Pavlova, M. Kambadur, M. Lewis, M. Si, M. K. Singh, M. Hassan, N. Goyal, N. Torabi, N. Bashlykov, N. Bogoychev, N. Chatterji, N. Zhang, O. Duchenne, O. Çelebi, P. Alrassy, P. Zhang, P. Li, P. Vasic, P. Weng, P. Bhargava, P. Dubal, P. Krishnan, P. S. Koura, P. Xu, Q. He, Q. Dong, R. Srinivasan, R. Ganapathy, R. Calderer, R. S. Cabral, R. Stojnic, R. Raileanu, R. Maheswari, R. Girdhar, R. Patel, R. Sauvestre, R. Polidoro, R. Sumbaly, R. Taylor, R. Silva, R. Hou, R. Wang, S. Hosseini, S. Chennabasappa, S. Singh, S. Bell, S. S. Kim, S. Edunov, S. Nie, S. Narang, S. R. Parthi, S. Shen, S. Wan, S. Bhosale, S. Zhang, S. Vandenheide, S. Batra, S. Whitman, S. Sootla, S. Collet, S. Gururangan, S. Borodinsky, T. Herman, T. Fowler, T. Sheasha, T. Georgiou, T. Scialom, T. Speckbacher, T. Mihaylov, T. Xiao, U. Karn, V. Goswami, V. Gupta, V. Ramanathan, V. Kerkez, V. Gonguet, V. Do, V. Vogeti, V. Albiero, V. Petrovic, W. Chu, W. Xiong, W. Fu, W. Meers, X. Martinet, X. Wang, X. Wang, X. E. Tan, X. Xia, X. Xie, X. Jia, X. Wang, Y. Goldschlag, Y. Gaur, Y. Babaei, Y. Wen, Y. Song, Y. Zhang, Y. Li,

References II

Y. Mao, Z. D. Coudert, Z. Yan, Z. Chen, Z. Papakipos, A. Singh, A. Srivastava, A. Jain, A. Kelsey, A. Shajnfeld, A. Gangidi, A. Victoria, A. Goldstand, A. Menon, A. Sharma, A. Boesenberg, A. Baevski, A. Feinstein, A. Kallet, A. Sangani, A. Teo, A. Yunus, A. Lupu, A. Alvarado, A. Caples, A. Gu, A. Ho, A. Poulton, A. Ryan, A. Ramchandani, A. Dong, A. Franco, A. Goyal, A. Saraf, A. Chowdhury, A. Gabriel, A. Bharambe, A. Eisenman, A. Yazdan, B. James, B. Maurer, B. Leonhardi, B. Huang, B. Loyd, B. D. Paola, B. Paranjape, B. Liu, B. Wu, B. Ni, B. Hancock, B. Wasti, B. Spence, B. Stojkovic, B. Gamido, B. Montalvo, C. Parker, C. Burton, C. Mejia, C. Liu, C. Wang, C. Kim, C. Zhou, C. Hu, C.-H. Chu, C. Cai, C. Tindal, C. Feichtenhofer, C. Gao, D. Civin, D. Beaty, D. Kreymer, D. Li, D. Adkins, D. Xu, D. Testuggine, D. David, D. Parikh, D. Liskovich, D. Foss, D. Wang, D. Le, D. Holland, E. Dowling, E. Jamil, E. Montgomery, E. Presani, E. Hahn, E. Wood, E.-T. Le, E. Brinkman, E. Arcaute, E. Dunbar, E. Smothers, F. Sun, F. Kreuk, F. Tian, F. Kokkinos, F. Ozgenel, F. Caggioni, F. Kanayet, F. Seide, G. M. Florez, G. Schwarz, G. Badeer, G. Swee, G. Halpern, G. Herman, G. Sizov, Guangyi, Zhang, G. Lakshminarayanan, H. Inan, H. Shojanazeri, H. Zou, H. Wang, H. Zha, H. Habeeb, H. Rudolph, H. Suk, H. Aspegren, H. Goldman, H. Zhan, I. Damlaj, I. Molybog, I. Tufanov, I. Leontiadis, I.-E. Veliche, I. Gat, J. Weissman, J. Geboski, J. Kohli, J. Lam, J. Asher, J.-B. Gaya, J. Marcus, J. Tang, J. Chan, J. Zhen, J. Reizenstein, J. Teboul, J. Zhong, J. Jin, J. Yang, J. Cummings, J. Carvill, J. Shepard, J. McPhie, J. Torres, J. Ginsburg, J. Wang, K. Wu, K. H. U, K. Saxena, K. Khandelwal, K. Zand, K. Matosich, K. Veeraraghavan, K. Michelena, K. Li, K. Jagadeesh, K. Huang, K. Chawla, K. Huang, L. Chen, L. Garg, L. A. L. Silva, L. Bell, L. Zhang, L. Guo, L. Yu, L. Moshkovich, L. Wehrstedt, M. Khabsa, M. Avalani, M. Bhatt, M. Mankus, M. Hasson, M. Lennie, M. Reso, M. Groshev, M. Naumov, M. Lathi, M. Keneally, M. Liu, M. L. Seltzer, M. Valko, M. Restrepo, M. Patel, M. Vyatskov, M. Samvelyan, M. Clark, M. Macey, M. Wang, M. J. Hermoso, M. Metanat, M. Rastegari, M. Bansal, N. Santhanam, N. Parks, N. White, N. Bawa, N. Singhal, N. Egebo, N. Usunier, N. Mehta, N. P. Laptev, N. Dong, N. Cheng, O. Chernoguz, O. Hart, O. Salpekar, O. Kalinli, P. Kent, P. Parekh, P. Saab, P. Balaji, P. Rittner, P. Bontrager, P. Roux, P. Dollar, P. Zvyagina, P. Ratanchandani, P. Yuvraj, Q. Liang, R. Alao, R. Rodriguez, R. Ayub, R. Murthy, R. Nayani, R. Mitra, R. Parthasarathy, R. Li, R. Hogan, R. Battey, R. Wang, R. Howes, R. Rinott, S. Mehta, S. Siby, S. J. Bondu, S. Datta, S. Chugh, S. Hunt, S. Dhillon, S. Sidorov, S. Pan, S. Mahajan, S. Verma, S. Yamamoto, S. Ramaswamy, S. Lindsay, S. Lindsay, S. Feng, S. Lin, S. C. Zha, S. Patil, S. Shankar, S. Zhang, S. Zhang, S. Wang, S. Agarwal, S. Sajuyigbe, S. Chintala, S. Max, S. Chen, S. Kehoe, S. Satterfield, S. Govindaprasad, S. Gupta, S. Deng, S. Cho, S. Virk, S. Subramanian, S. Choudhury, S. Goldman, T. Remez, T. Glaser, T. Best, T. Koehler, T. Robinson, T. Li, T. Zhang, T. Matthews, T. Chou, T. Shaked, V. Vontimitta,

References III

V. Ajayi, V. Montanez, V. Mohan, V. S. Kumar, V. Mangla, V. Ionescu, V. Poenaru, V. T. Mihailescu, V. Ivanov, W. Li, W. Wang, W. Jiang, W. Bouaziz, W. Constable, X. Tang, X. Wu, X. Wang, X. Wu, X. Gao, Y. Kleinman, Y. Chen, Y. Hu, Y. Jia, Y. Qi, Y. Li, Y. Zhang, Y. Zhang, Y. Adi, Y. Nam, Yu, Wang, Y. Zhao, Y. Hao, Y. Qian, Y. Li, Y. He, Z. Rait, Z. DeVito, Z. Rosnbrick, Z. Wen, Z. Yang, Z. Zhao, and Z. Ma, "The Llama 3 Herd of Models," Nov. 2024.

- [5] OpenAI, A. Jaech, A. Kalai, A. Lerer, A. Richardson, A. El-Kishky, A. Low, A. Helyar, A. Madry, A. Beutel, A. Carney, A. Iftimie, A. Karpenko, A. T. Passos, A. Neitz, A. Prokofiev, A. Wei, A. Tam, A. Bennett, A. Kumar, A. Saraiva, A. Vallone, A. Duberstein, A. Kondrich, A. Mishchenko, A. Applebaum, A. Jiang, A. Nair, B. Zoph, B. Ghorbani, B. Zhang, B. Rossen, B. Sokolowsky, B. Barak, B. McGrew, B. Minaiev, B. Hao, B. Baker, B. Houghton, B. McKinzie, B. Eastman, C. Lugaresi, C. Bassin, C. Hudson, C. M. Li, C. de Bourcy, C. Voss, C. Shen, C. Zhang, C. Koch, C. Orsinger, C. Hesse, C. Fischer, C. Chan, D. Roberts, D. Kappler, D. Levy, D. Selsam, D. Dohan, D. Farhi, D. Mely, D. Robinson, D. Tsipras, D. Li, D. Oprica, E. Freeman, E. Zhang, E. Wong, E. Proehl, E. Cheung, E. Mitchell, E. Wallace, E. Ritter, E. Mays, F. Wang, F. P. Such, F. Raso, F. Leoni, F. Tsimplouras, F. Song, F. von Lohmann, F. Sulit, G. Salmon, G. Parascandolo, G. Chabot, G. Zhao, G. Brockman, G. Leclerc, H. Salman, H. Bao, H. Sheng, H. Andrin, H. Bagherinezhad, H. Ren, H. Lightman, H. W. Chung, I. Kivlichan, I. O'Connell, I. Osband, I. C. Gilaberte, I. Akkaya, I. Kostrikov, I. Sutskever, I. Kofman, J. Pachocki, J. Lennon, J. Wei, J. Harb, J. Twore, J. Feng, J. Yu, J. Weng, J. Tang, J. Yu, J. Q. Candela, J. Palermo, J. Parish, J. Heidecke, J. Hallman, J. Rizzo, J. Gordon, J. Uesato, J. Ward, J. Huizinga, J. Wang, K. Chen, K. Xiao, K. Singhal, K. Nguyen, K. Cobbe, K. Shi, K. Wood, K. Rimbach, K. Gu-Lemberg, K. Liu, K. Lu, K. Stone, K. Yu, L. Ahmad, L. Yang, L. Liu, L. Maksin, L. Ho, L. Fedus, L. Weng, L. Li, L. McCallum, L. Held, L. Kuhn, L. Kondraciuk, L. Kaiser, L. Metz, M. Boyd, M. Trebacz, M. Joglekar, M. Chen, M. Tintor, M. Meyer, M. Jones, M. Kaufer, M. Schwarzer, M. Shah, M. Yatbaz, M. Y. Guan, M. Xu, M. Yan, M. Glaese, M. Chen, M. Lampe, M. Malek, M. Wang, M. Fradin, M. McClay, M. Pavlov, M. Wang, M. Wang, M. Murati, M. Bavarian, M. Rohaninejad, N. McAleese, N. Chowdhury, N. Chowdhury, N. Ryder, N. Tezak, N. Brown, O. Nachum, O. Boiko, O. Murk, O. Watkins, P. Chao, P. Ashbourne, P. Izmailov, P. Zhokhov, R. Dias, R. Arora, R. Lin, R. G. Lopes, R. Gaon, R. Miyara, R. Leike, R. Hwang, R. Garg, R. Brown, R. James, R. Shu, R. Cheu, R. Greene, S. Jain, S. Altman, S. Toizer, S. Toyer, S. Miserendino, S. Agarwal, S. Hernandez, S. Baker, S. McKinney, S. Yan, S. Zhao, S. Hu, S. Santurkar, S. R. Chaudhuri, S. Zhang, S. Fu, S. Papay, S. Lin, S. Balaji, S. Sanjeev, S. Sidor, T. Broda, A. Clark, T. Wang, T. Gordon, T. Sanders, T. Patwardhan, T. Sottiaux, T. Degry, T. Dimson, T. Zheng, T. Garipov, T. Stasi, T. Bansal, T. Creech, T. Peterson, T. Eloundou, V. Qi,

References IV

- V. Kosaraju, V. Monaco, V. Pong, V. Fomenko, W. Zheng, W. Zhou, W. Zhan, W. McCabe, W. Zaremba, Y. Dubois, Y. Lu, Y. Chen, Y. Cha, Y. Bai, Y. He, Y. Zhang, Y. Wang, Z. Shao, and Z. Li, "OpenAI o1 System Card," Apr. 2026.
- [6] A. Jake, "Uber burned through its entire 2026 AI budget in four months. Now its COO is questioning whether it's worth it."
- [7] Q. Cheng, P. Lin, and Y. Zhao, "Does generative AI facilitate investor Trading? Early evidence from ChatGPT outages," *Journal of Accounting and Economics*, vol. 80, p. 101821, Nov. 2025.
- [8] Z. Wang, J. Wohlwend, and T. Lei, "Structured Pruning of Large Language Models," in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, (Online), pp. 6151–6162, Association for Computational Linguistics, 2020.
- [9] G. Hinton, O. Vinyals, and J. Dean, "Distilling the Knowledge in a Neural Network," Mar. 2015.
- [10] J. Wu, Y. Wang, Z. Wu, Z. Wang, A. Veeraraghavan, and Y. Lin, "Deep \$k\$-Means: Re-Training and Parameter Sharing with Harder Cluster Assignments for Compressing Deep Convolutions," June 2018.
- [11] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference," Dec. 2017.
- [12] C. E. Shannon, "A mathematical theory of communication," *The Bell System Technical Journal*, vol. 27, no. 3, pp. 379–423, 1948.
- [13] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, C. Zheng, D. Liu, F. Zhou, F. Huang, F. Hu, H. Ge, H. Wei, H. Lin, J. Tang, J. Yang, J. Tu, J. Zhang, J. Yang, J. Yang, J. Zhou, J. Zhou, J. Lin, K. Dang, K. Bao, K. Yang, L. Yu, L. Deng, M. Li, M. Xue, M. Li, P. Zhang, P. Wang, Q. Zhu, R. Men, R. Gao, S. Liu, S. Luo, T. Li, T. Tang, W. Yin, X. Ren, X. Wang, X. Zhang, X. Ren, Y. Fan, Y. Su, Y. Zhang, Y. Zhang, Y. Wan, Y. Liu, Z. Wang, Z. Cui, Z. Zhang, Z. Zhou, and Z. Qiu, "Qwen3 Technical Report," May 2025.
- [14] Qwen Team, "Qwen3.5: Accelerating productivity with native multimodal agents," Feb. 2026.

References V

- [15] Gemma Team, S. E. Abd, V. Aggarwal, R. Algayres, A. Andreev, O. Bachem, I. Ballantyne, C. Brick, V. Cărbune, M. Casbon, M. Chaturvedi, A. Chawla, V. Cotruta, A. Coucke, P. Culliton, R. Dadashi, L. Dixon, M. Elhawaty, U. Evci, C. Farabet, J. Ferret, F. Galgani, S. Girgin, J.-B. Grill, M. Grootendorst, J. Guo, C. Hardin, Y. He, S. M. Hernandez, O. Homburger, L. Hussenot, J. Ji, A. Joulin, A. Kamath, P. Kassraie, O. Lacombe, P. Lahoti, G. Liu, G. Martins, L. Martins, T. Matejovicova, R. Merhej, N. Momchev, S. Mondal, R. Mullins, S. R. Panyam, S. Pathak, S. Perrin, A. S. Pinto, E. Pot, A. Pouget, A. Ramé, S. Ramos, D. Reid, D. Rim, M. Rivière, K. Roth, L. Rouillard, O. Sanseviero, P. G. Sessa, S. Settle, D. Sinopalnikov, S. Smoot, P. Stanczyk, A. Steiner, L. Stewart, I. Tolstikhin, M. Tschannen, A. Tsitsulin, N. Vieillard, R. Wu, P. Xu, H. Yang, E. Yvinec, B. Zhang, L. Zhang, J. Zou, N. Aagnes, A. Abdelhamed, J. Adamek, S. Agrawal, S. Agrawal, I. Alabdulmohsin, J. B. Alayrac, U. Alon, C. Amarnath, A. Anand, C. Anastasiou, S. Ariafar, F.-X. Aubet, K. Axiotis, F. Barbero, J. Barral, A. Bendebury, U. Bergmann, S. Bileschi, K. Black, M. Blondel, S. Borgeaud, A. Bražiškas, R. Burnell, R. Busa-Fekete, M. Cai, D. Calandriello, G. Cameron, C. Caucheteux, R. Chaabouni, G. Chadha, J. Chan, B. J. Chen, J. Chen, L. Chen, X. Chen, D. Cheng, T.-h. Chien, N. Chinaev, Y. Chou, Z. Chu, B. Coleman, P. Consul, S. Conway-Rahman, S. Crowell, D. Cutler, V. Dani, S. Daruki, A. Das, D. Deutsch, N. Dikkala, L. Ding, Q. Ding, S. Dodhia, K. Donhauser, T. Doshi, A. Dragan, A. Druinsky, S. Dua, Z. Egyed, D. Eisenbud, D. Eppens, C. Fan, B. Fatemi, Y. Fathullah, V. Feinberg, M. Ferev, S. Flennerhag, T. Fujimoto, J. G. Oliveira, I. Galatzer-Levy, J. Gante, S. Geisler, S. Ghosal, A. M. Girgis, T. von Glehn, A. Go, A. Gokhale, A. Grills, Y. Gu, M. Gupta, P. Gupta, G. Guruganesh, R. Hadsell, H. Harkous, J. Harlalka, D. Hassabis, A. Hauth, J. Heyward, A. Hosseini, C.-Y. Hsia, I.-H. Hsu, X. Huang, Y. Huang, K. Hui, A. Hutter, T. I, F. Iliopoulos, A. Jain, G. Jawahar, Z. Ji, Q. Jin, M. Johnson, K. Joshi, A. Kandoor, W.-C. Kang, K. Kavukcuoglu, M. Kazemi, K. Kenealy, A. Khalifa, P. Kirk, I. Korotkov, S. Kothawade, V. Kovalev, N. Kovelamudi, A. Kraft, R. Kumar, V. Kumar, H. Kuppam, J. Lannin, C.-Y. Lee, S. Lee, D. Lepikhin, A. Levkovitch, D. Li, Q. Li, V. Liévin, E. Lin, Z. Lin, C. Liu, T. Liu, T. Liu, X. Liu, I. Lobov, M. Lunayach, M. Ma, G. Madan, A. Maksai, E. Malmi, M. Matuszak, D. McDuff, G. Menghani, M. Miśkula, D. Mirylenka, K. Misiunas, V. Misra, A. Mitran, K. Mohamed, M. Mukha, E. Noland, J. O'Donnell, B. O'Donoghue, K. Olszewska, B. Orlando, W. Pan, R. Panigrahy, U. Parekh, N. Perez-Nieves, C. Park, E. Paskie, L. Peng, B. Petrini, S. Petrov, J. Pfeiffer, B. Piot, M. Plomecka, S. Poder, O. Ponce, A. Pramanik, D. Racz, A. Rajan, M. Ramanovich, A. Rao, M. Ritter, V. Rodrigues, E. Rosen, M. Rybiński, N. Sachdeva, M. E. Sander, R. Sathyanarayana, S. Savla, S. Schmidgall, T. Schuster, G. Scrivener, B. Seguin, A. Sellergrén, A. Severyn, I. Shafran, D. Shah, B. Shahriari, Y. Shangguan, A. Shenoy, P. Shenoy, R. Shivanna, P. Sho, L. Spangher, W. Stokowiec, T. Strother, Y. Su, Y. Sun,

References VI

- M. Sundararajan, A. Tacchetti, M. H. Taege, P. Tafti, J. Tarbouriech, C. Tekur, S. Thakoor, R. Thapa, M. Traverse, L. Treven, T. Tu, C. T. Tung, Ç. Ünlü, P. Veličković, M. P. Venkat, S. G. Venkatesh, V. Venkiteswaran, F. Visin, A. Vitvitskyi, K. Vodrahalli, W. Wang, X. Wang, T. Warkentin, J. Wassenberg, J. Wieting, C. Wu, L. Xiao, H. Xu, Y. Xu, F. Xue, A. Yadav, J. Yan, A. Yang, L. Yang, M.-H. Yang, Z. Ying, J. H. Yoo, M. Zadimoghaddam, S. Zafar, F. Zhang, J. Zhang, J. Zhang, X. Zhang, C. Zhao, D. Zhou, and C. Zou, "Gemma 4 Technical Report," July 2026.
- [16] Z.ai inc., "GLM-4.7: Advancing the Coding Capability," Dec. 2025.
- [17] R. Eldan and Y. Li, "TinyStories: How Small Can Language Models Be and Still Speak Coherent English?," May 2023.
- [18] S. Black, L. Gao, P. Wang, C. Leahy, and S. Biderman, "GPT-Neo: Large scale autoregressive language modeling with meshtensorflow." Zenodo, Oct. 2021.
- [19] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, "Language models are unsupervised multitask learners," 2019.
- [20] T. Kudo and J. Richardson, "SentencePiece: A simple and language independent subword tokenizer and detokenizer for Neural Text Processing," Aug. 2018.
- [21] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, A. Rodriguez, A. Joulin, E. Grave, and G. Lample, "LLaMA: Open and Efficient Foundation Language Models," Feb. 2023.
- [22] F. Jelinek, R. L. Mercer, L. R. Bahl, and J. K. Baker, "Perplexity—a measure of the difficulty of speech recognition tasks," *The Journal of the Acoustical Society of America*, vol. 62, pp. S63–S63, Dec. 1977.
- [23] M. Nagel, R. A. Amjad, M. van Baalen, C. Louizos, and T. Blankevoort, "Up or Down? Adaptive Rounding for Post-Training Quantization," June 2020.
- [24] J. M. Cohen, S. Kaur, Y. Li, J. Z. Kolter, and A. Talwalkar, "Gradient Descent on Neural Networks Typically Occurs at the Edge of Stability," Nov. 2022.

References VII

- [25] Kimi Team, Y. Bai, Y. Bao, Y. Charles, C. Chen, G. Chen, H. Chen, H. Chen, J. Chen, N. Chen, R. Chen, Y. Chen, Y. Chen, Y. Chen, Z. Chen, J. Cui, H. Ding, M. Dong, A. Du, C. Du, D. Du, Y. Du, Y. Fan, Y. Feng, K. Fu, B. Gao, C. Gao, H. Gao, P. Gao, T. Gao, Y. Ge, S. Geng, Q. Gu, X. Gu, L. Guan, H. Guo, J. Guo, X. Hao, T. He, W. He, W. He, Y. He, C. Hong, H. Hu, Y. Hu, Z. Hu, W. Huang, Z. Huang, Z. Huang, T. Jiang, Z. Jiang, X. Jin, Y. Kang, G. Lai, C. Li, F. Li, H. Li, M. Li, W. Li, Y. Li, Y. Li, Y. Li, Z. Li, Z. Li, H. Lin, X. Lin, Z. Lin, C. Liu, C. Liu, H. Liu, J. Liu, J. Liu, L. Liu, S. Liu, T. Y. Liu, T. Liu, W. Liu, Y. Liu, Y. Liu, Y. Liu, Z. Liu, E. Lu, H. Lu, L. Lu, Y. Luo, S. Ma, X. Ma, Y. Ma, S. Mao, J. Mei, X. Men, Y. Miao, S. Pan, Y. Peng, R. Qin, Z. Qin, B. Qu, Z. Shang, L. Shi, S. Shi, F. Song, J. Su, Z. Su, L. Sui, X. Sun, F. Sung, Y. Tai, H. Tang, J. Tao, Q. Teng, C. Tian, C. Wang, D. Wang, F. Wang, H. Wang, H. Wang, J. Wang, J. Wang, J. Wang, S. Wang, S. Wang, S. Wang, X. Wang, Y. Wang, Y. Wang, Y. Wang, Y. Wang, Y. Wang, Z. Wang, Z. Wang, Z. Wang, Z. Wang, C. Wei, Q. Wei, H. Wu, W. Wu, X. Wu, Y. Wu, C. Xiao, J. Xie, X. Xie, W. Xiong, B. Xu, J. Xu, L. H. Xu, L. Xu, S. Xu, W. Xu, X. Xu, Y. Xu, Z. Xu, J. Xu, J. Xu, J. Yan, Y. Yan, H. Yang, X. Yang, Y. Yang, Y. Yang, Z. Yang, Z. Yang, Z. Yang, H. Yao, X. Yao, W. Ye, Z. Ye, B. Yin, L. Yu, E. Yuan, H. Yuan, M. Yuan, S. Yuan, H. Zhan, D. Zhang, H. Zhang, W. Zhang, X. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Y. Zhang, Z. Zhang, H. Zhao, Y. Zhao, Z. Zhao, H. Zheng, S. Zheng, L. Zhong, J. Zhou, X. Zhou, Z. Zhou, J. Zhu, Z. Zhu, W. Zhuang, and X. Zu, "Kimi K2: Open Agentic Intelligence," Feb. 2026.
- [26] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, "Scikit-learn: Machine learning in python," *Journal of Machine Learning Research*, vol. 12, no. 85, pp. 2825–2830, 2011.
- [27] E. Borenstein, E. Sharon, and S. Ullman, "Combining top-down and bottom-up segmentation," vol. 2004, pp. 46–46, July 2004.
- [28] D. S. Kermamy, M. Goldbaum, W. Cai, C. C. Valentim, H. Liang, S. L. Baxter, A. McKeown, G. Yang, X. Wu, F. Yan, J. Dong, M. K. Prasadha, J. Pei, M. Y. Ting, J. Zhu, C. Li, S. Hewett, J. Dong, I. Ziyar, A. Shi, R. Zhang, L. Zheng, R. Hou, W. Shi, X. Fu, Y. Duan, V. A. Huu, C. Wen, E. D. Zhang, C. L. Zhang, O. Li, X. Wang, M. A. Singer, X. Sun, J. Xu, A. Tafreshi, M. A. Lewis, H. Xia, and K. Zhang, "Identifying Medical Diagnoses and Treatable Diseases by Image-Based Deep Learning," *Cell*, vol. 172, pp. 1122–1131.e9, Feb. 2018.
- [29] I. Loshchilov and F. Hutter, "Decoupled Weight Decay Regularization," Jan. 2019.
- [30] I. Loshchilov and F. Hutter, "SGDR: Stochastic Gradient Descent with Warm Restarts," May 2017.